NoSQL

ENGR689 (Sprint)

How to Store Data In Cloud?

- Networked File Systems

Centralized server
with attached storage media

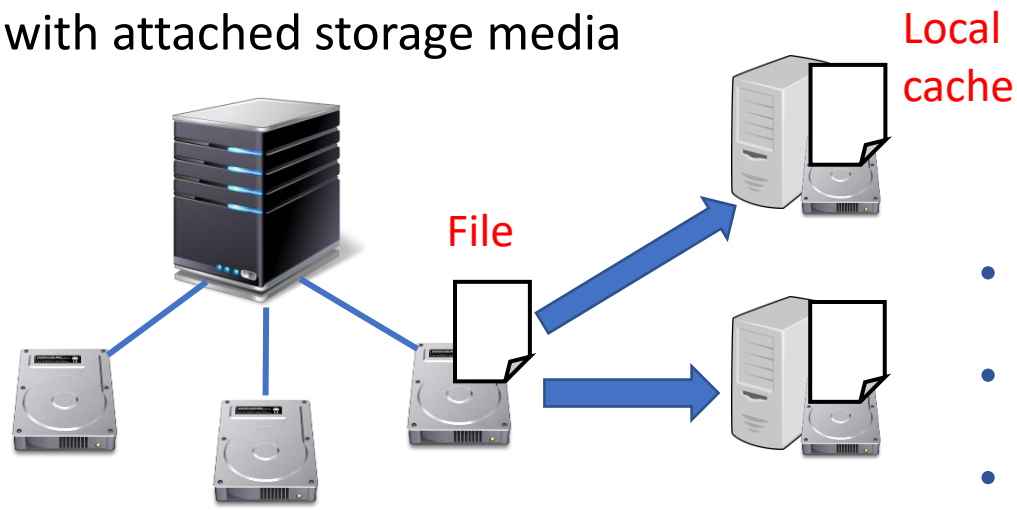

- Hierarchical structures
- Difficult to scale for more storage
- Good for static contents
- Extremely weak consistency (open-close consistency)

RDBMS
(Relational Database Management Systems)

- RDBMS has:
  - Predefined schemas with tuples / records / rows
  - Well-defined SQL interface

```
SELECT Name, SUM(Grade) FROM
students JOIN Records
ON Students.ID = Records.ID
GROUP BY Students.ID
```

- Easy to program join/union operations
- Strong consistency for updates

Schema - Students

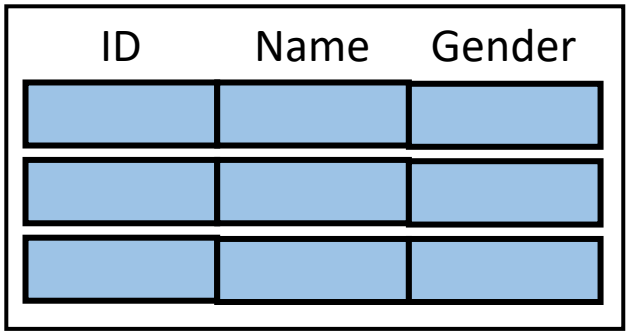

Schema - Records

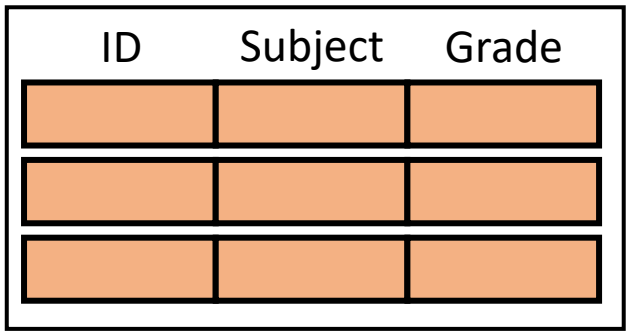

How to Scale RDBMS

- Exchanging data on partitioned DB is expensive

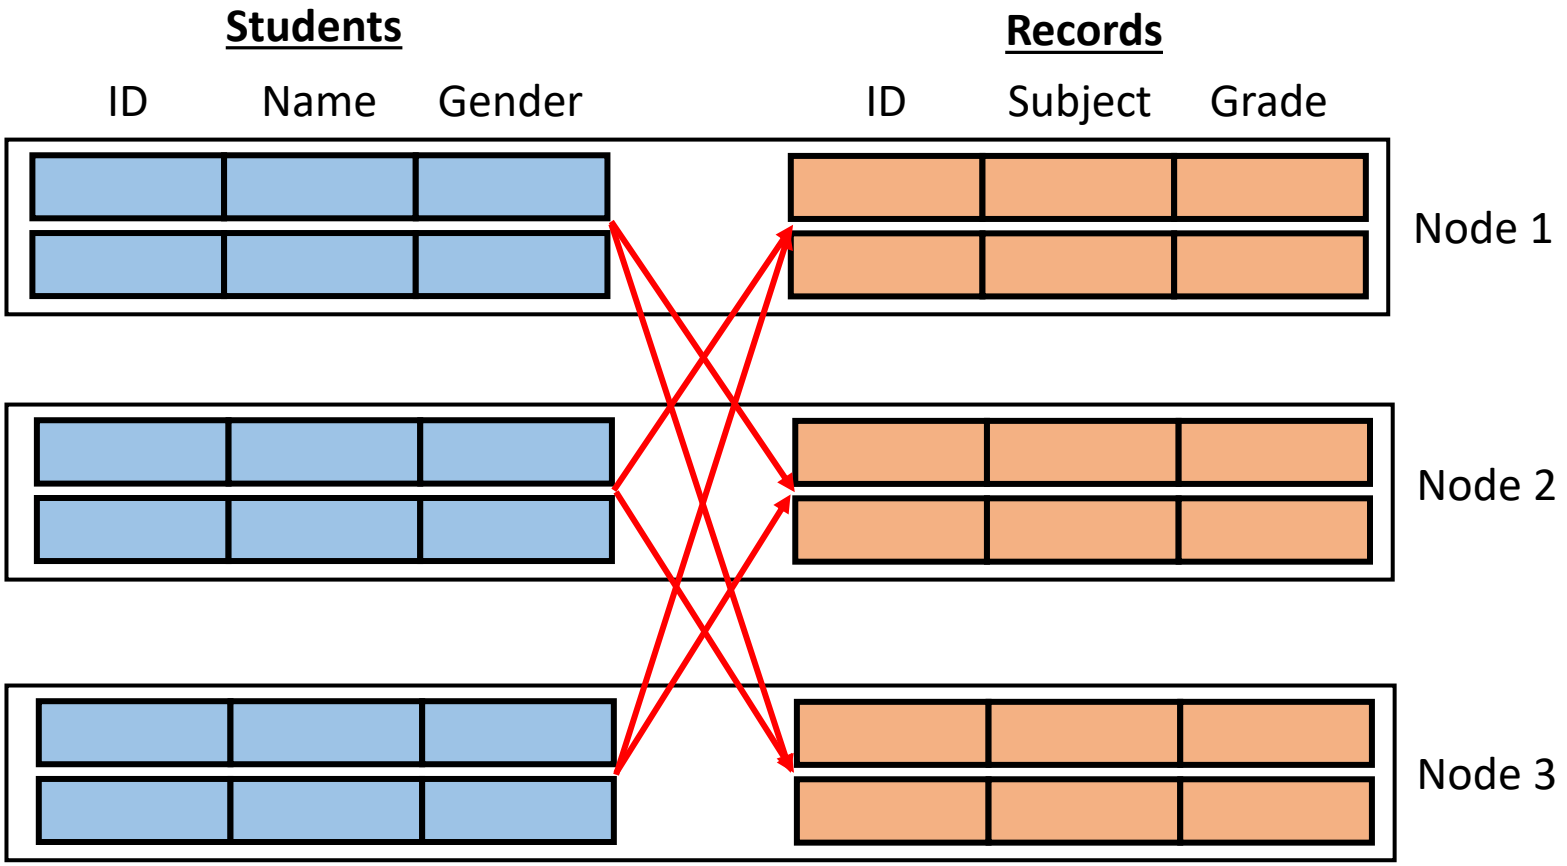

Data movement becomes the main bottleneck

Dynamo

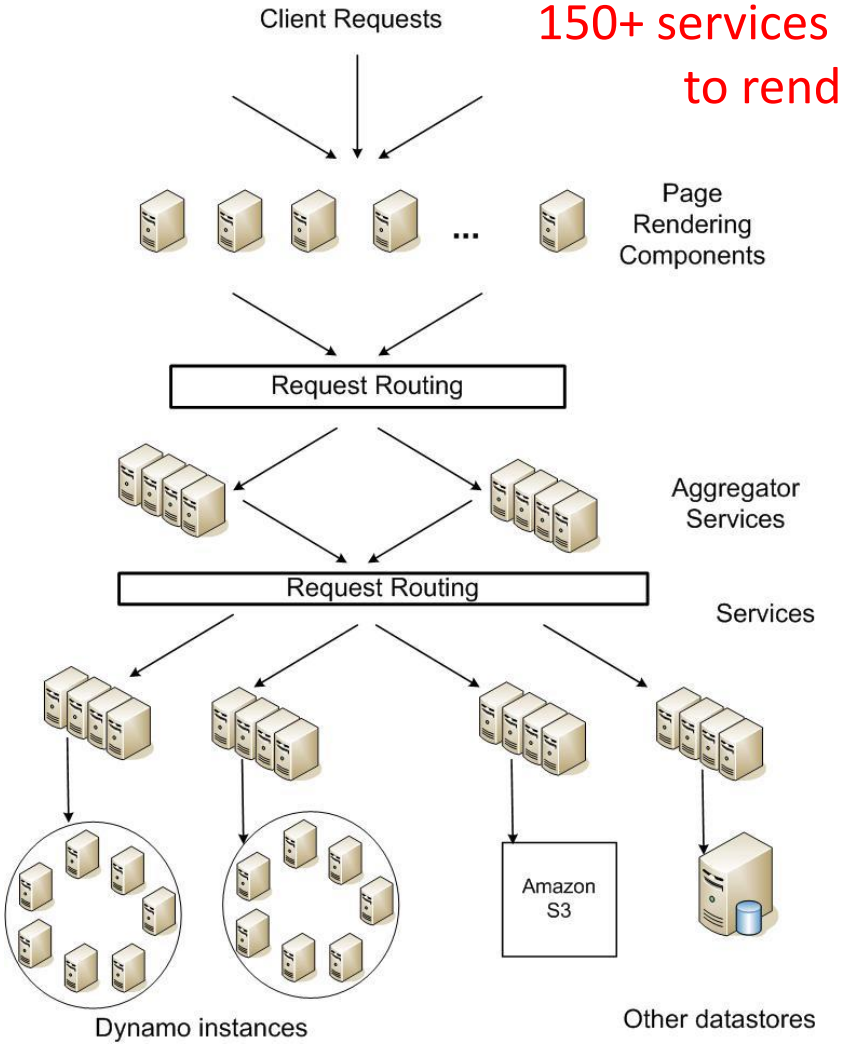

150+ services inter-connected to render a page

- Highly scalable and available key-value store for Amazon
- Service Level Agreement (SLA):
  99.9% of requests receive responses within 300ms

NoSQL (Not Only SQL)

- Auto-sharding without relational operations
- Optimized for simple queries (e.g., what's the name of UID 012345?)
- Map: Key →
  - Value (Key-Value Stores, e.g., Dynamo, Memcached)
  - Document (Document Stores, e.g., MongoDB)
  - Node with connections (Graph Stores, e.g., Neo4j)
  - Columns (Wide Column Stores, e.g., HBase, Cassandra)

Simple Key-Value Store API

- Get(Key) → Value
- Put(Key, Value)

High Availability: Both Get and Put need to finish immediately, how?

- Multiple replicas to ensure no blocking on node failure
- Tolerate temporary inconsistency to prevent blocking on synchronization

Partitioning & Replication

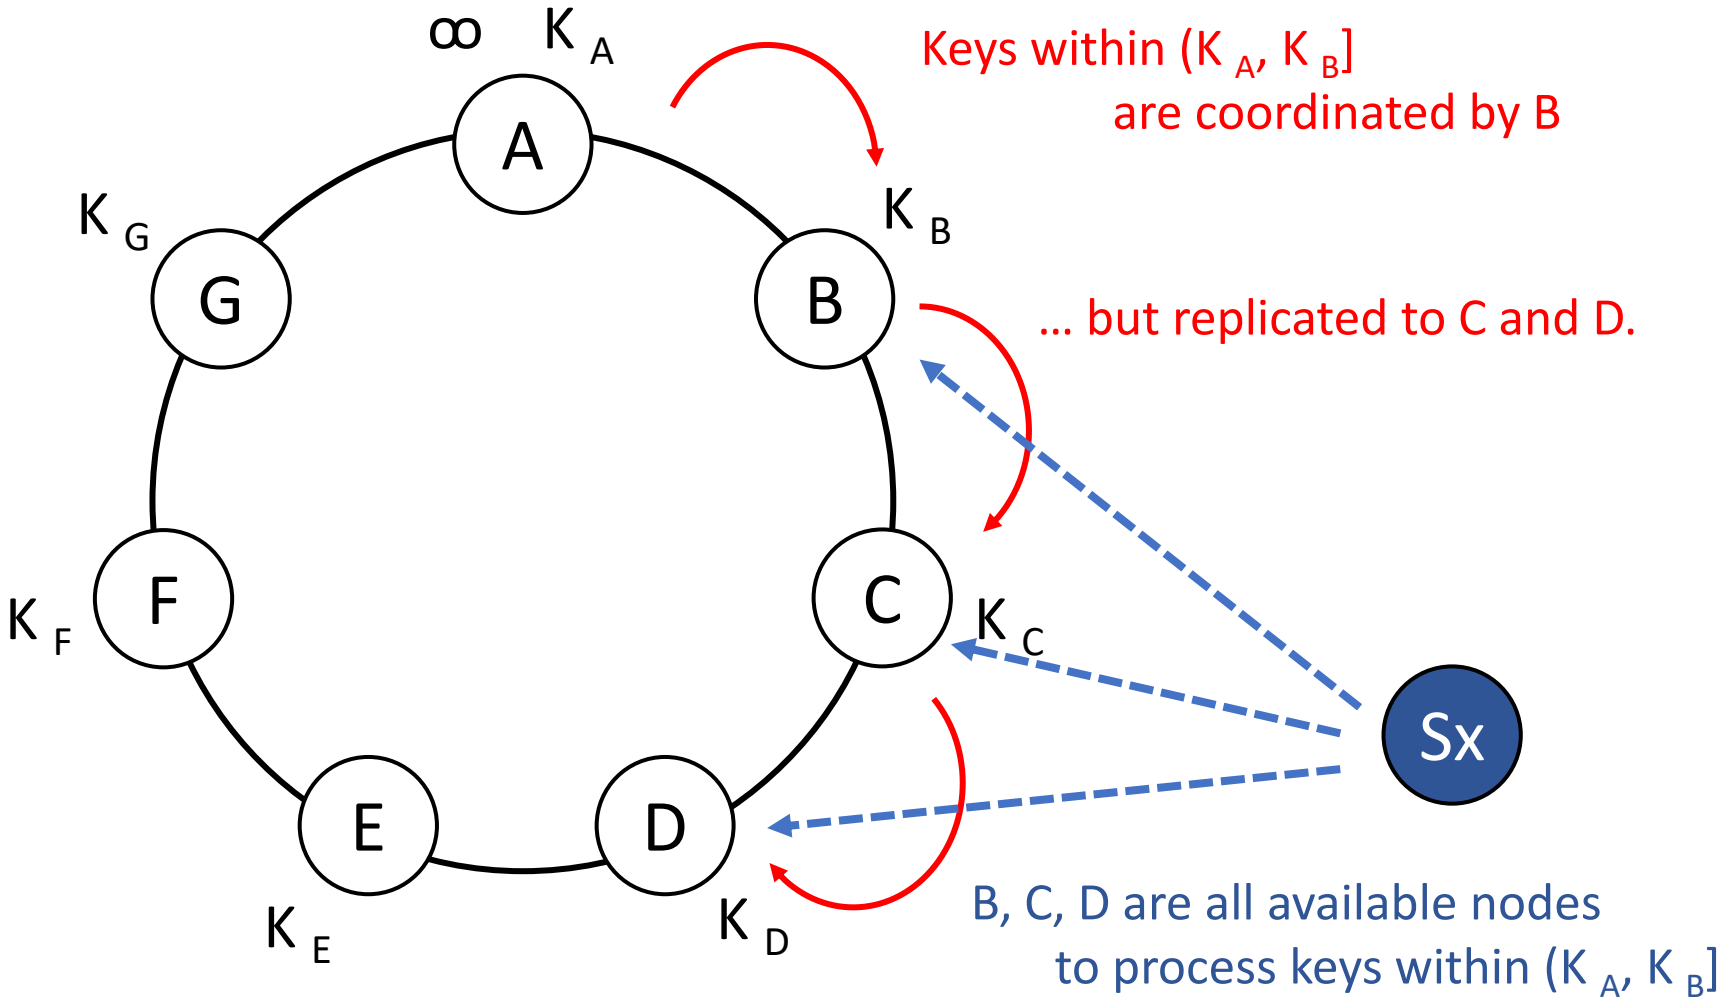

Eventual Consistency

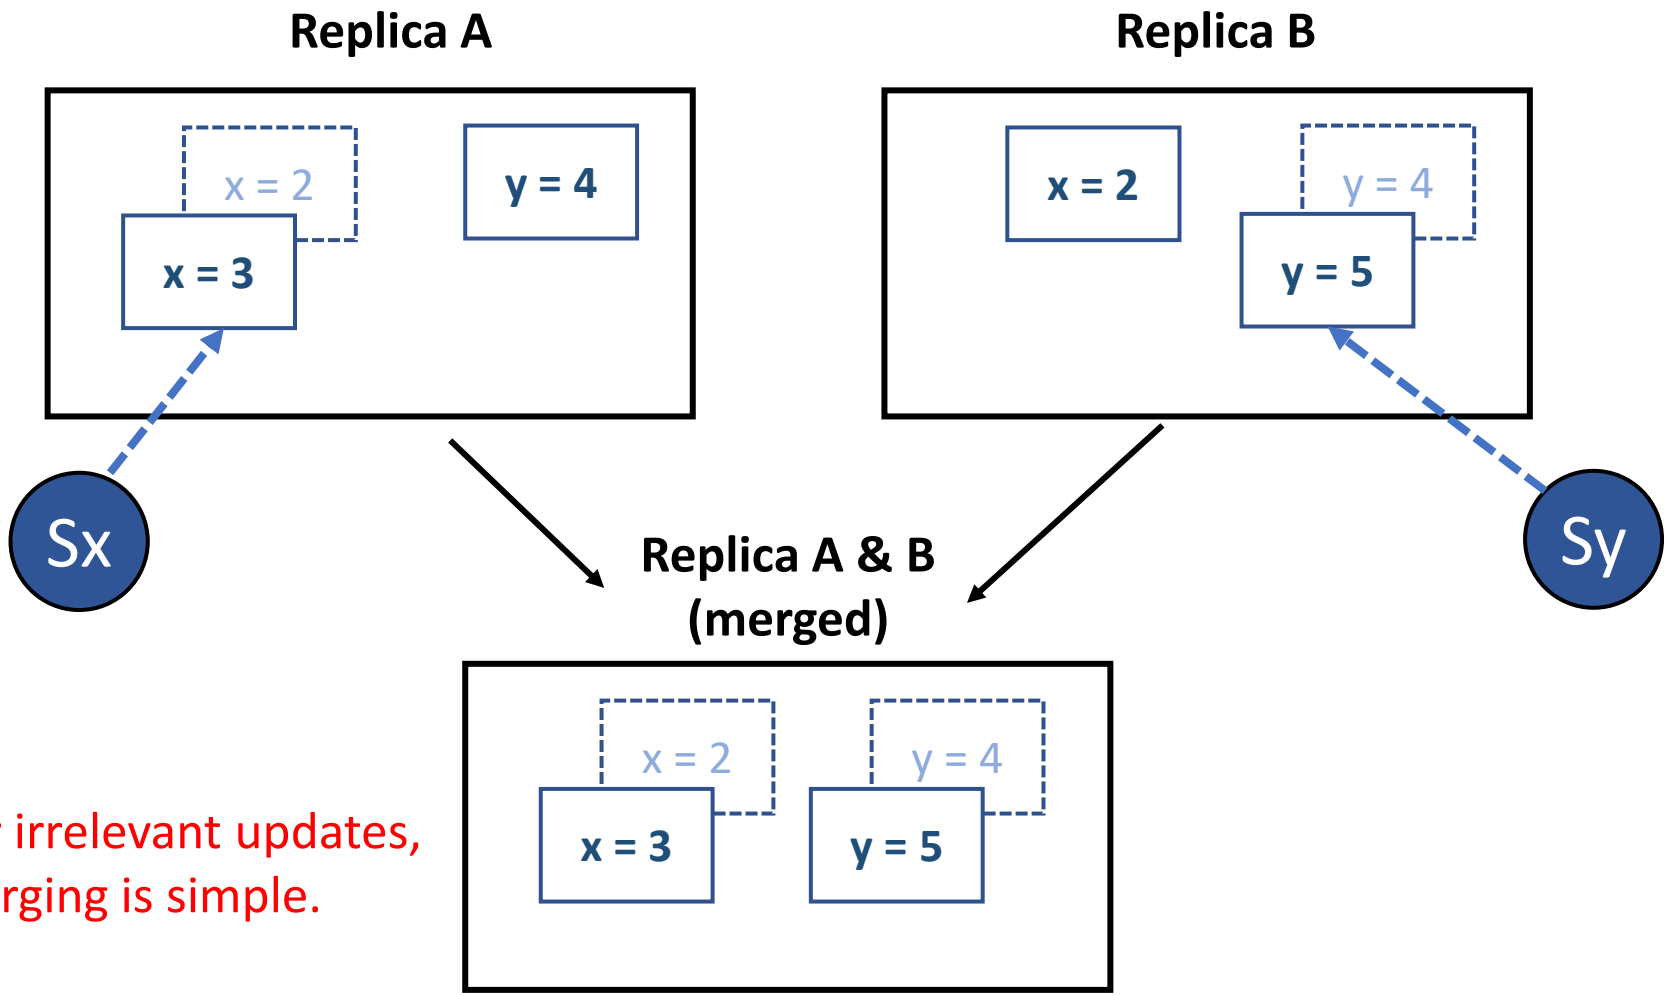

For irrelevant updates, merging is simple.

Eventual Consistency

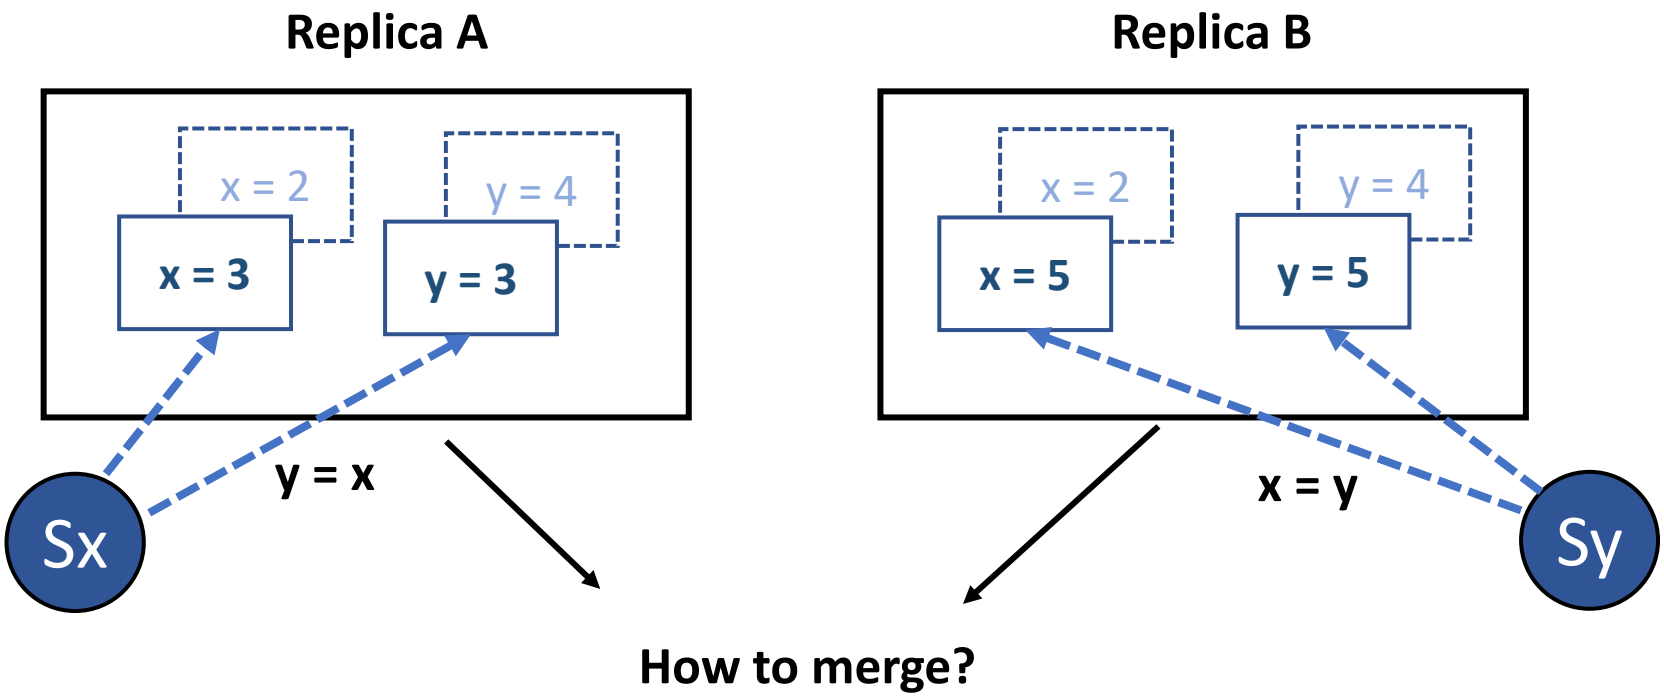

Need to keep track of the "causal" relationship:
i.e., knowing x and y are updated by the same users/servers

Versioned Data

- Keeping multiple versions until being resolved by a server

Vectored clock: a virtual clock advanced at every put() to keep track of data dependencies.

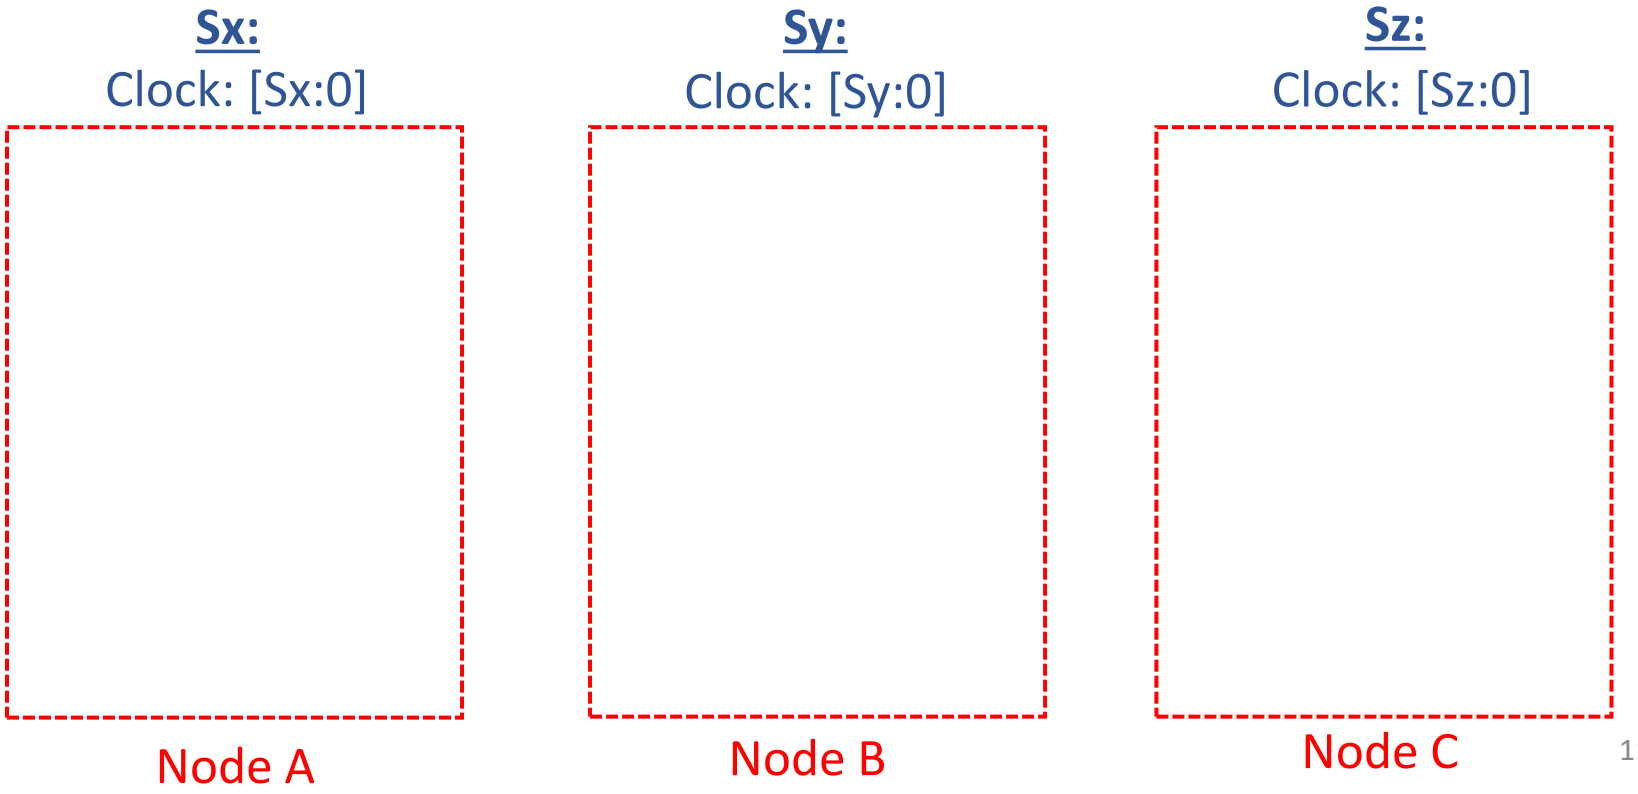

Versioned Data

- Keeping multiple versions until being resolved by a server

Vectored clock: a virtual clock advanced at every put() to keep track of data dependencies.

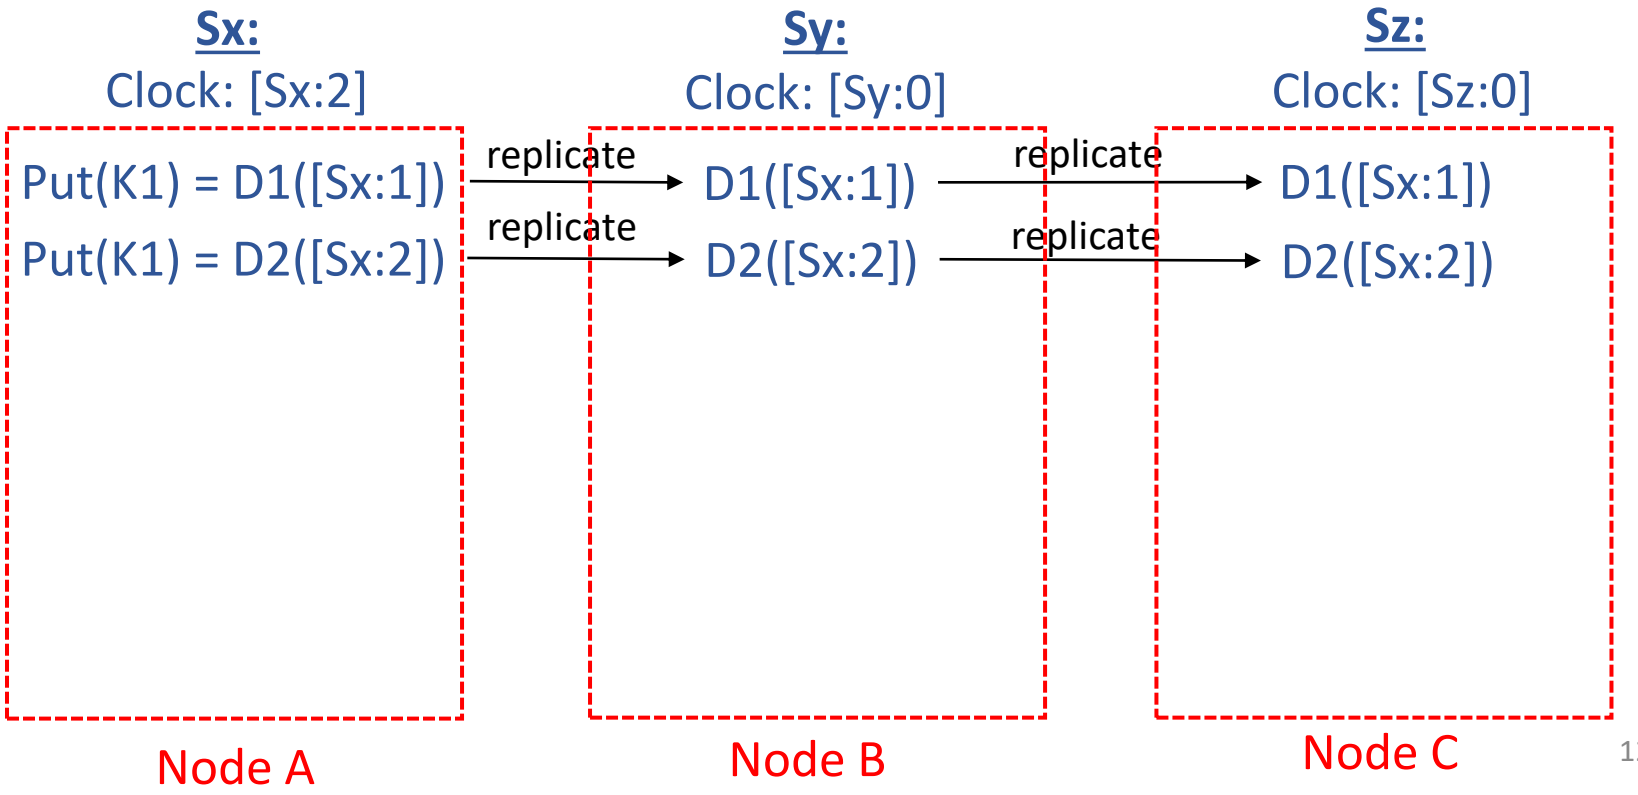

Versioned Data

- Keeping multiple versions until being resolved by a server

Vectored clock: a virtual clock advanced at every put() to keep track of data dependencies.

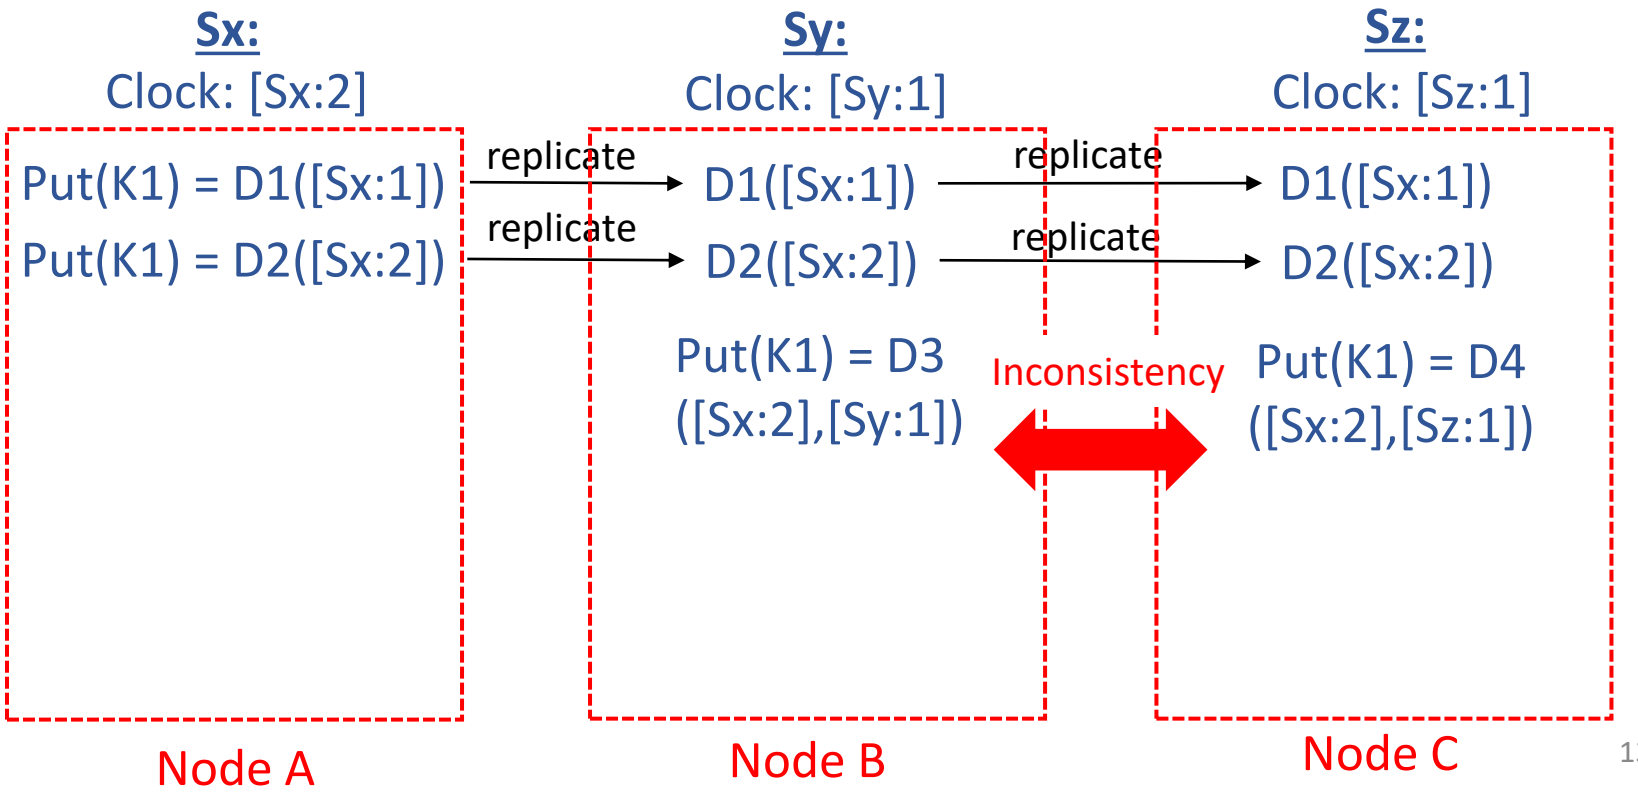

Versioned Data

- Keeping multiple versions until being resolved by a server

Vectored clock: a virtual clock advanced at every put() to keep track of data dependencies.

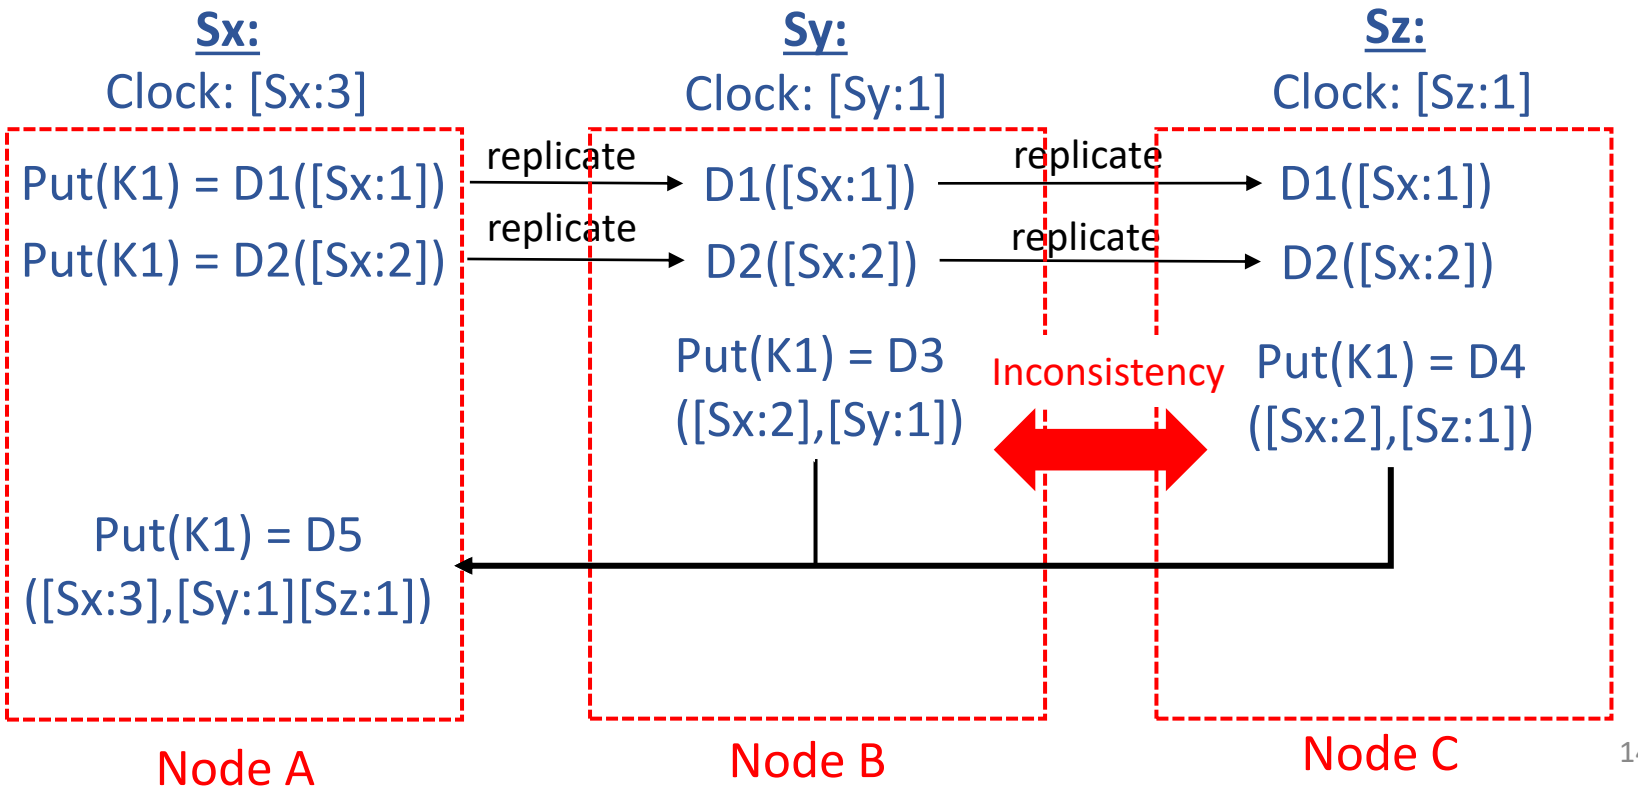

Versioned Data

- Keeping multiple versions until being resolved by a server

Vectored clock: a virtual clock advanced at every put() to keep track of data dependencies.

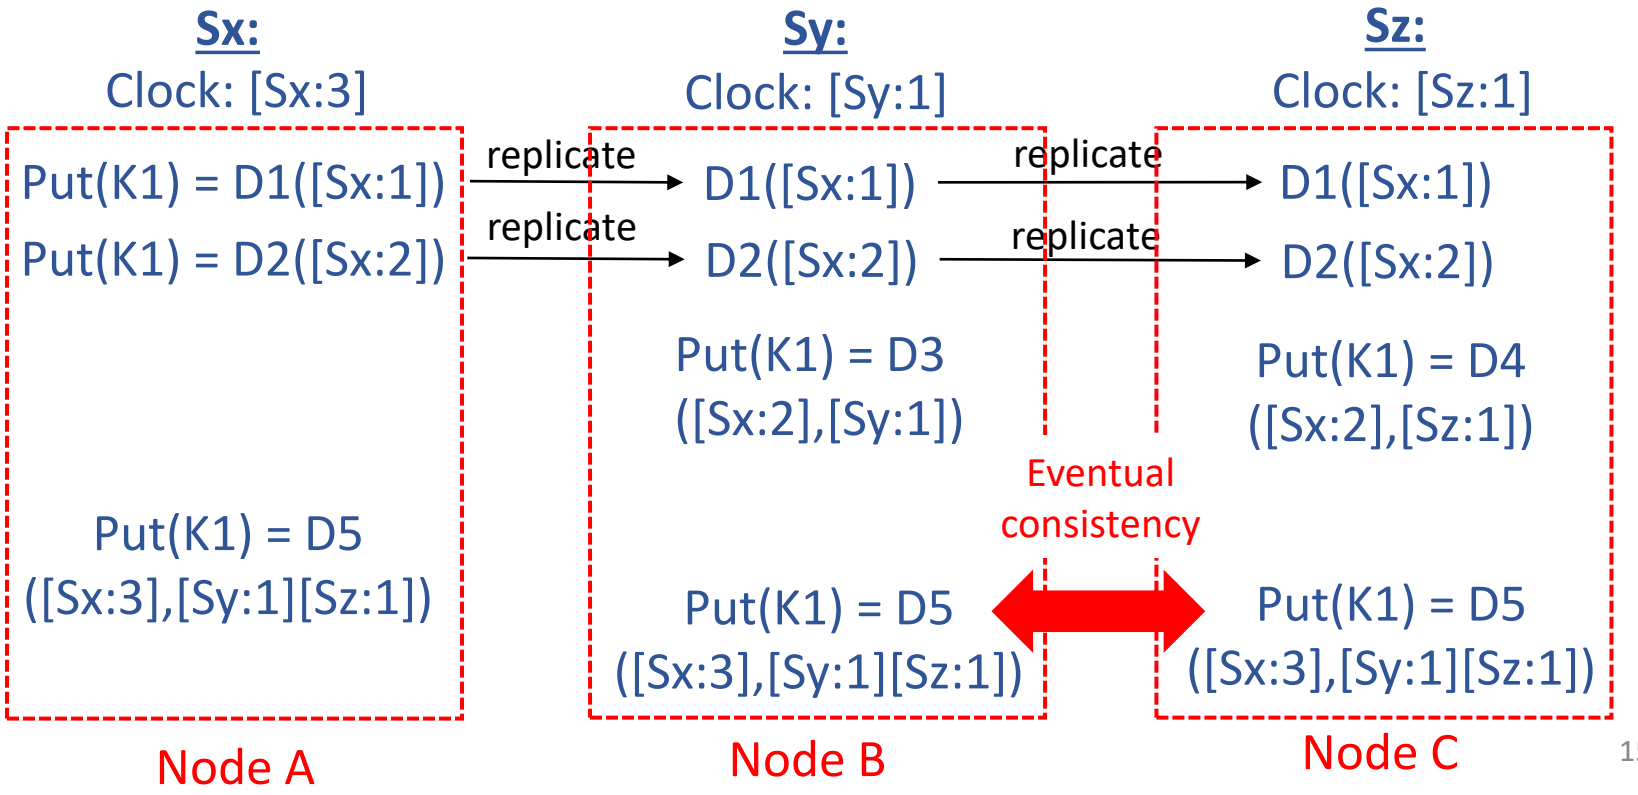

Dynamo API

- Get(key) → Value or [(Value, Context)]

  A reader either get() a consistent value,
  or a list of versioned values with respective contexts (clocks).

- Put(Key, Context, Value)

  A writer can decide which context to put() the new value.

Failure Detection

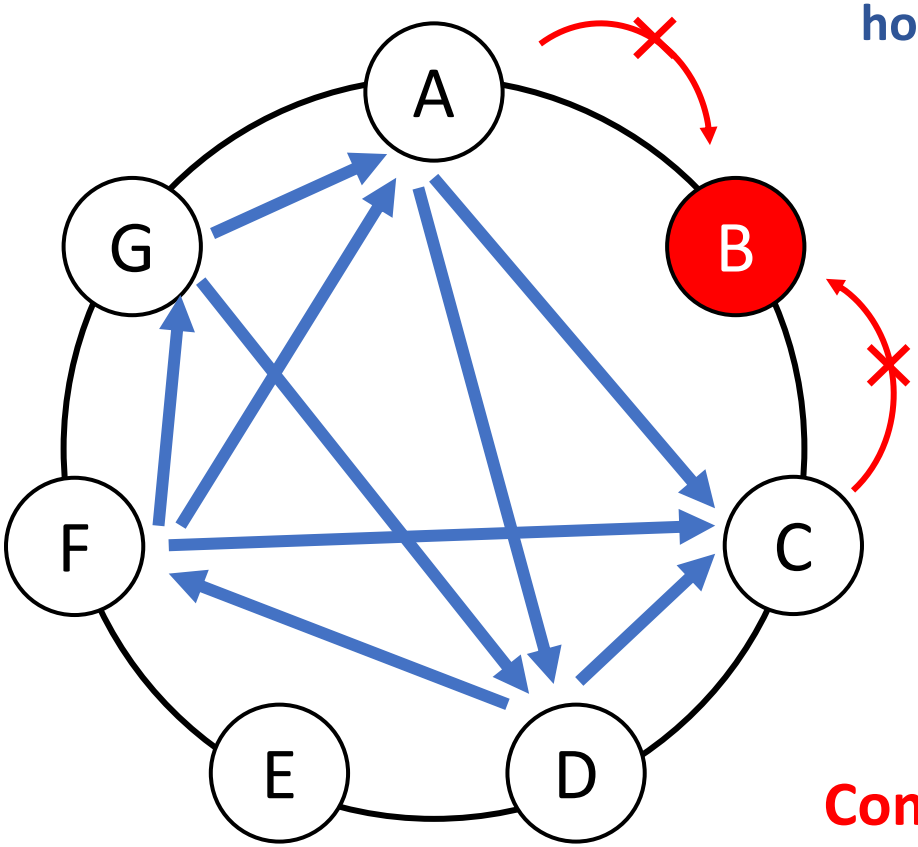

If a node is offline,
how do other nodes know about it?

Ans: gossip-based protocol

- Node A and C consider node B offline if receive no response from B
- Each node randomly contacts other node every second to exchange node information.

Completely decentralized failure detection protocol

Recover From Temporary Failures

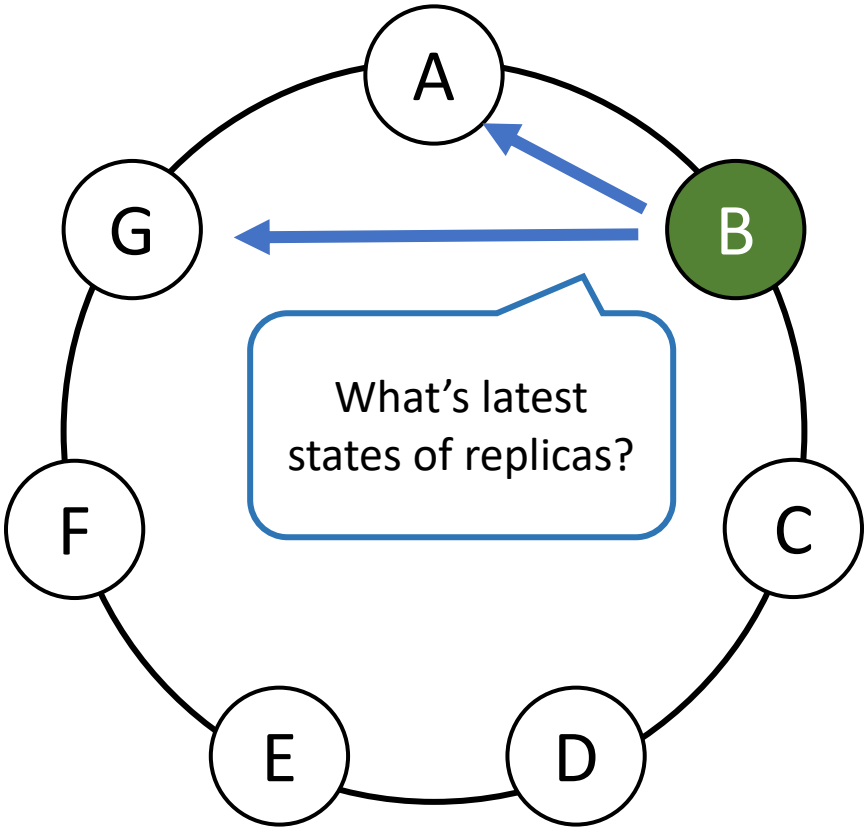

- If B becomes available again, it needs to synchronize with other replicas.
- Checking the state of replicas can be expensive
  - Sending all the data
    ➔ Too expensive
  - Sending hash of the replica
    ➔ Need rehashing at updates

Summary

- NoSQL simplifies both syntaxes and requirements of distributed storage
  - Open()/Read()/Write() or SQL queries ➔ Get() and Put()
- Dynamo uses classic distributed system techniques to ensure high availability (99.9 SLA)
  - Versioned data with vectored clocks
  - Gossip-based failure detection