


Virtualization

ENGR 689 (Sprint)




What's Virtualization?

Physical vs Virtual Machines



Without virtualization



With virtualization

Why Virtualization in Cloud?

- **Ease of resource sharing:**
splitting hardware resources for multiple tenants
- **Security Isolation:**
tenants are isolated from each other
- **Legacy software:**
tenants run their OSES and applications without modification
- **Flexibility of deployment:**
tenants can be easily deployed, migrated, or removed from a physical host

Classic Virtualization

Popek & Goldberg (1974):

- **Identical environment:**

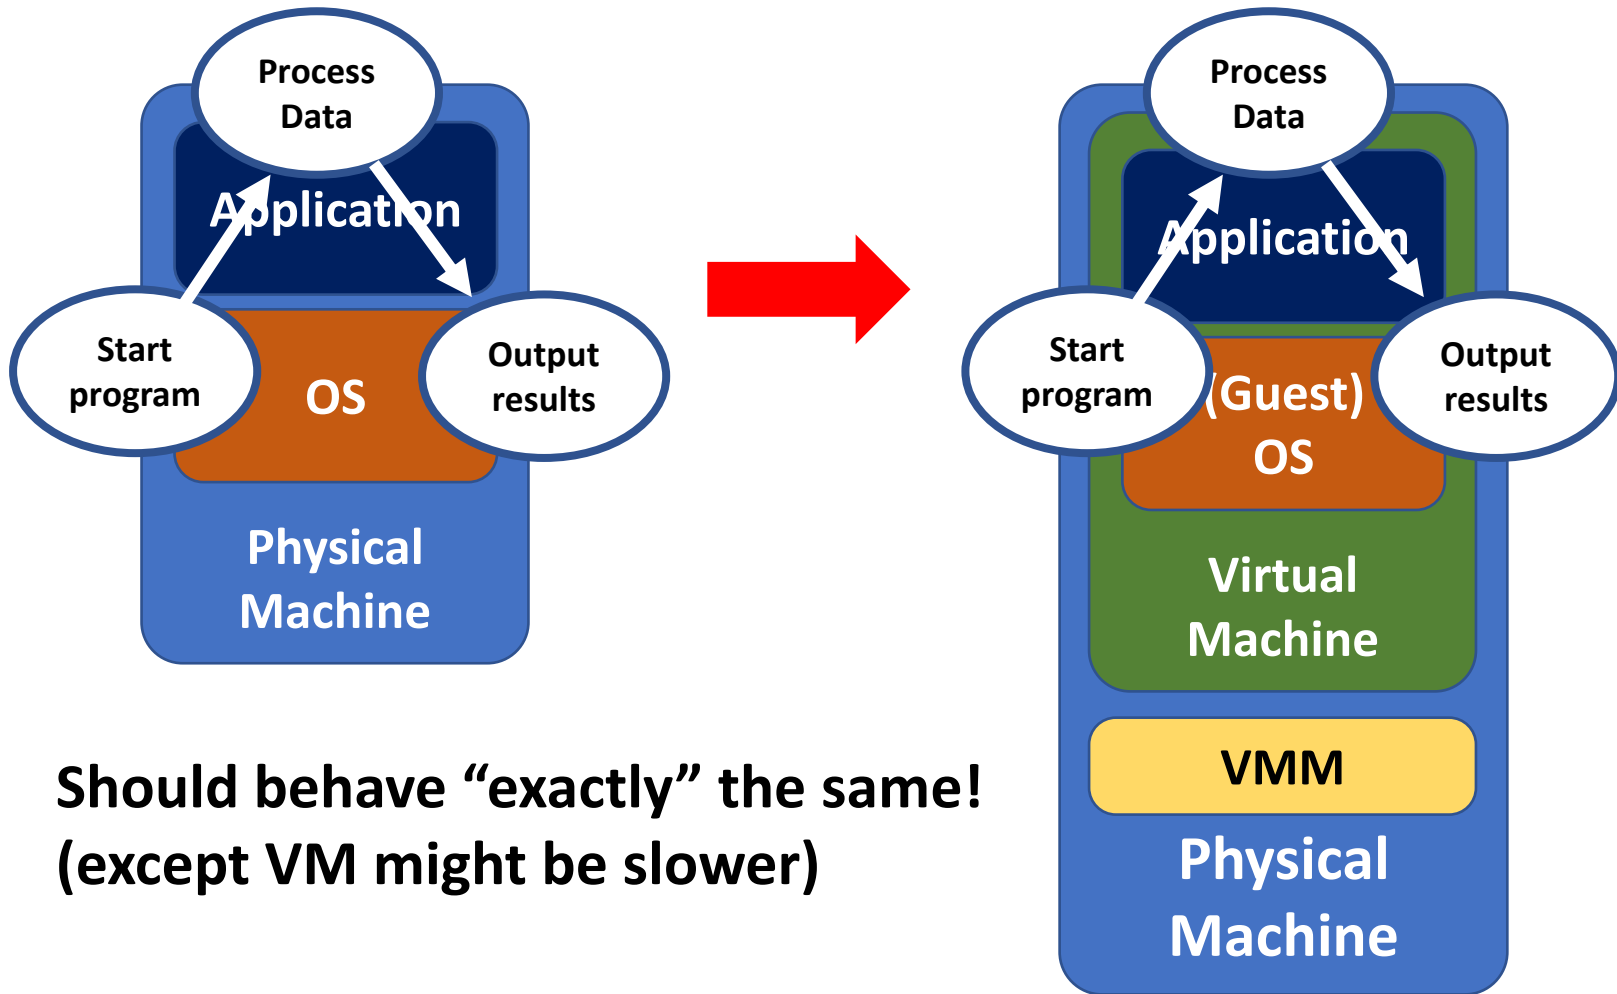
Programs in VMs should have identical effects as running on real machines

- **Efficiency:** should not add too much overheads

- **Resource control:**

Virtual machine monitor (VMM) has complete control of hardware resources

Identity



Should behave “exactly” the same!
(except VM might be slower)

How to Achieve Virtualization?

- Software emulation
- Hardware-based virtualization
- Paravirtualization (not classic virtualization)
 - Discuss in next lecture

Software Emulation

Guest OS Binary

00000000	push	ebp
00000001	mov	ebp, esp
00000003	movzx	ecx, [ebp+arg_0]
00000007	pop	ebp
00000008	movzx	dx, cl
0000000C	lea	eax, [edx+edx]
0000000F	add	eax, edx
00000011	shl	eax, 2
00000014	add	eax, edx
00000016	shr	eax, 8
00000019	sub	cl, al
0000001B	shr	cl, 1
0000001D	add	al, cl
0000001F	shr	al, 5
00000022	movzx	eax, al
00000025	retn	

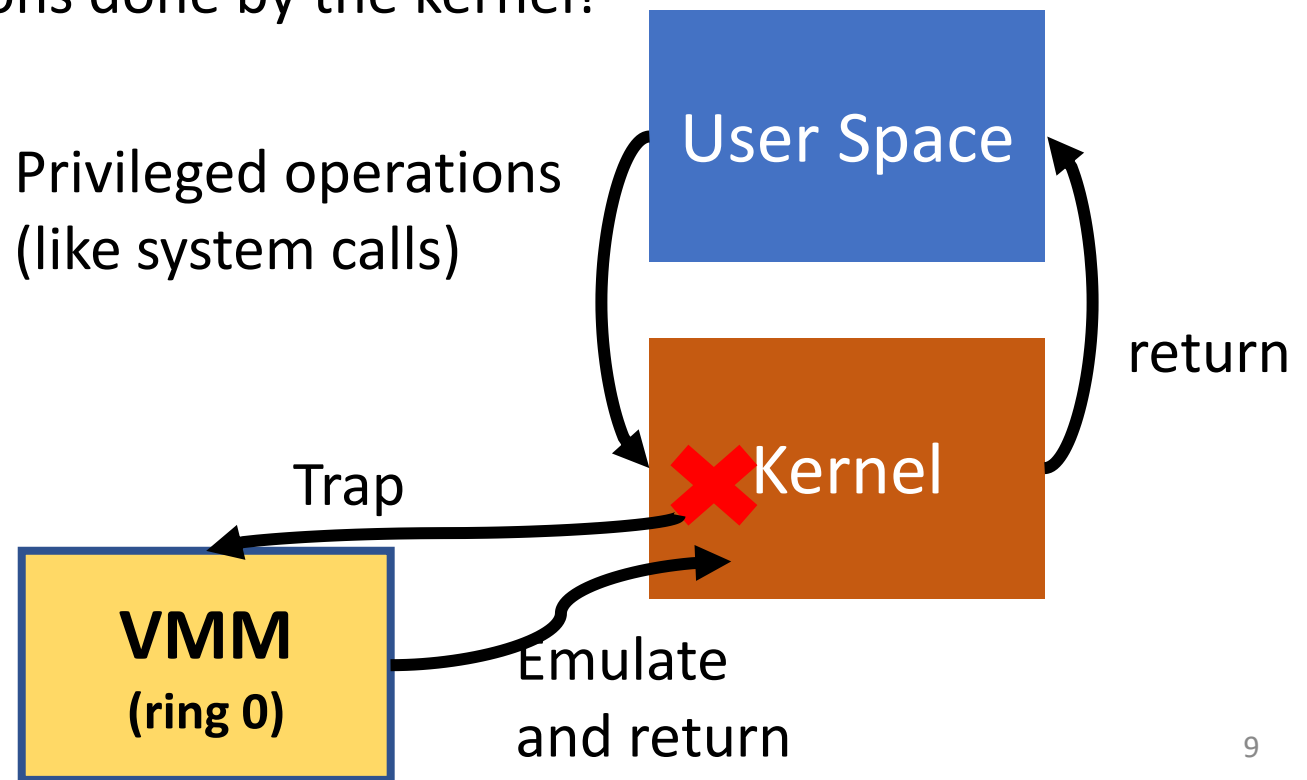
Emulate the execution
of instructions one by one.



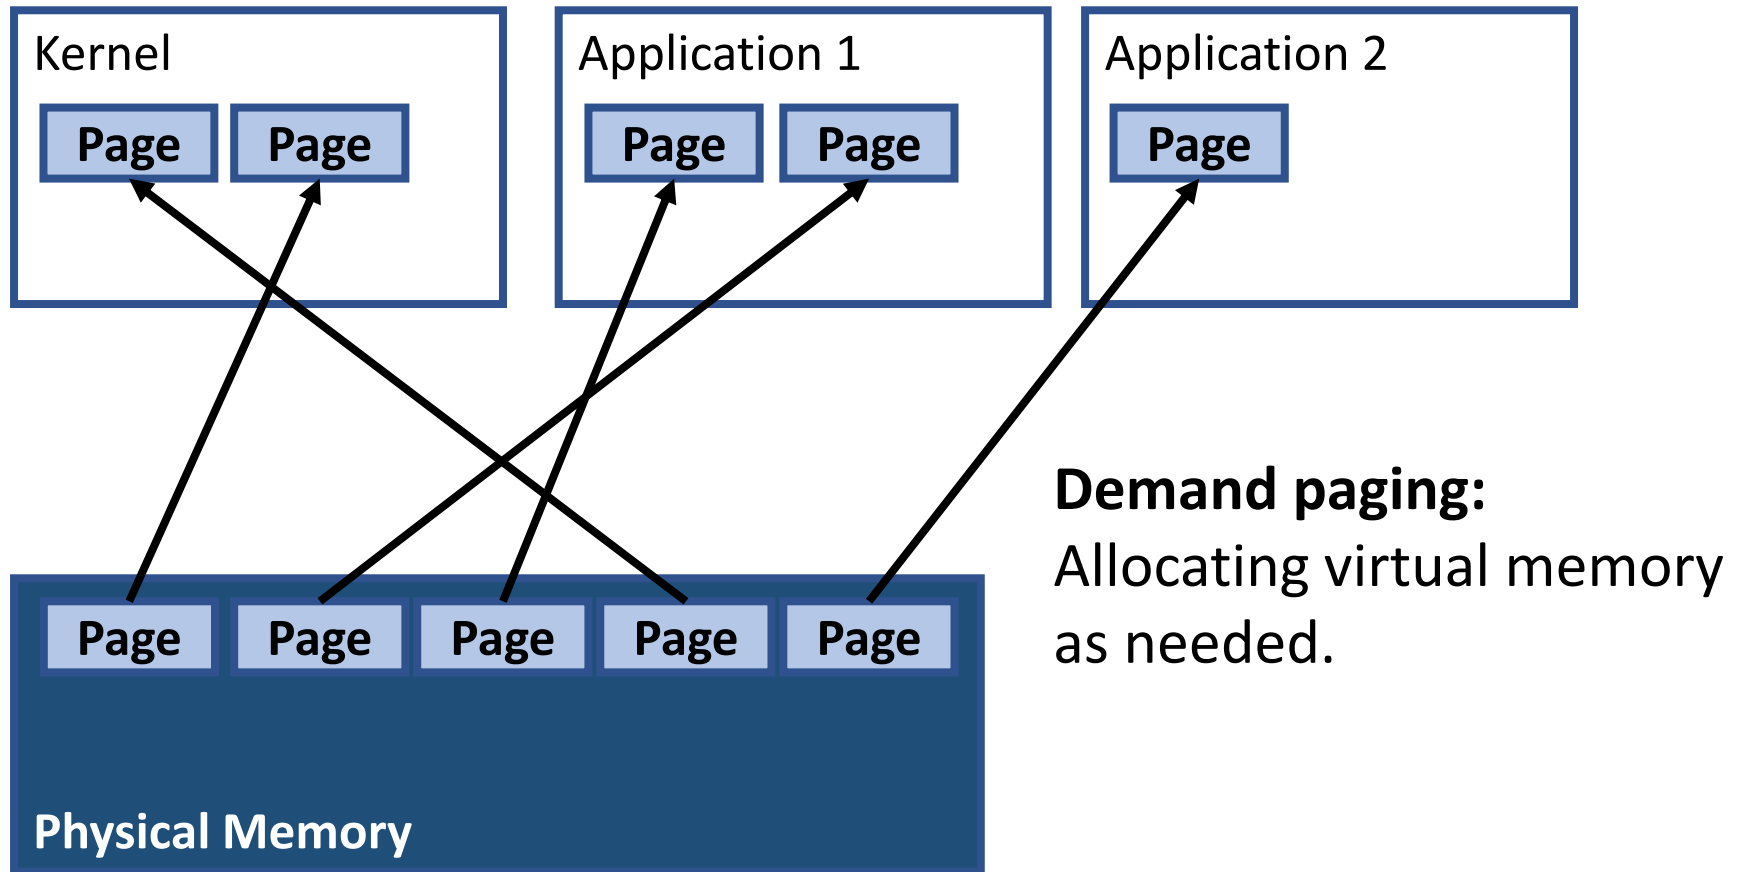
Emulator
(e.g., Qemu)

Trap & Emulate

- For efficiency, only “privileged operations” need to be emulated
 - Operations done by the kernel!



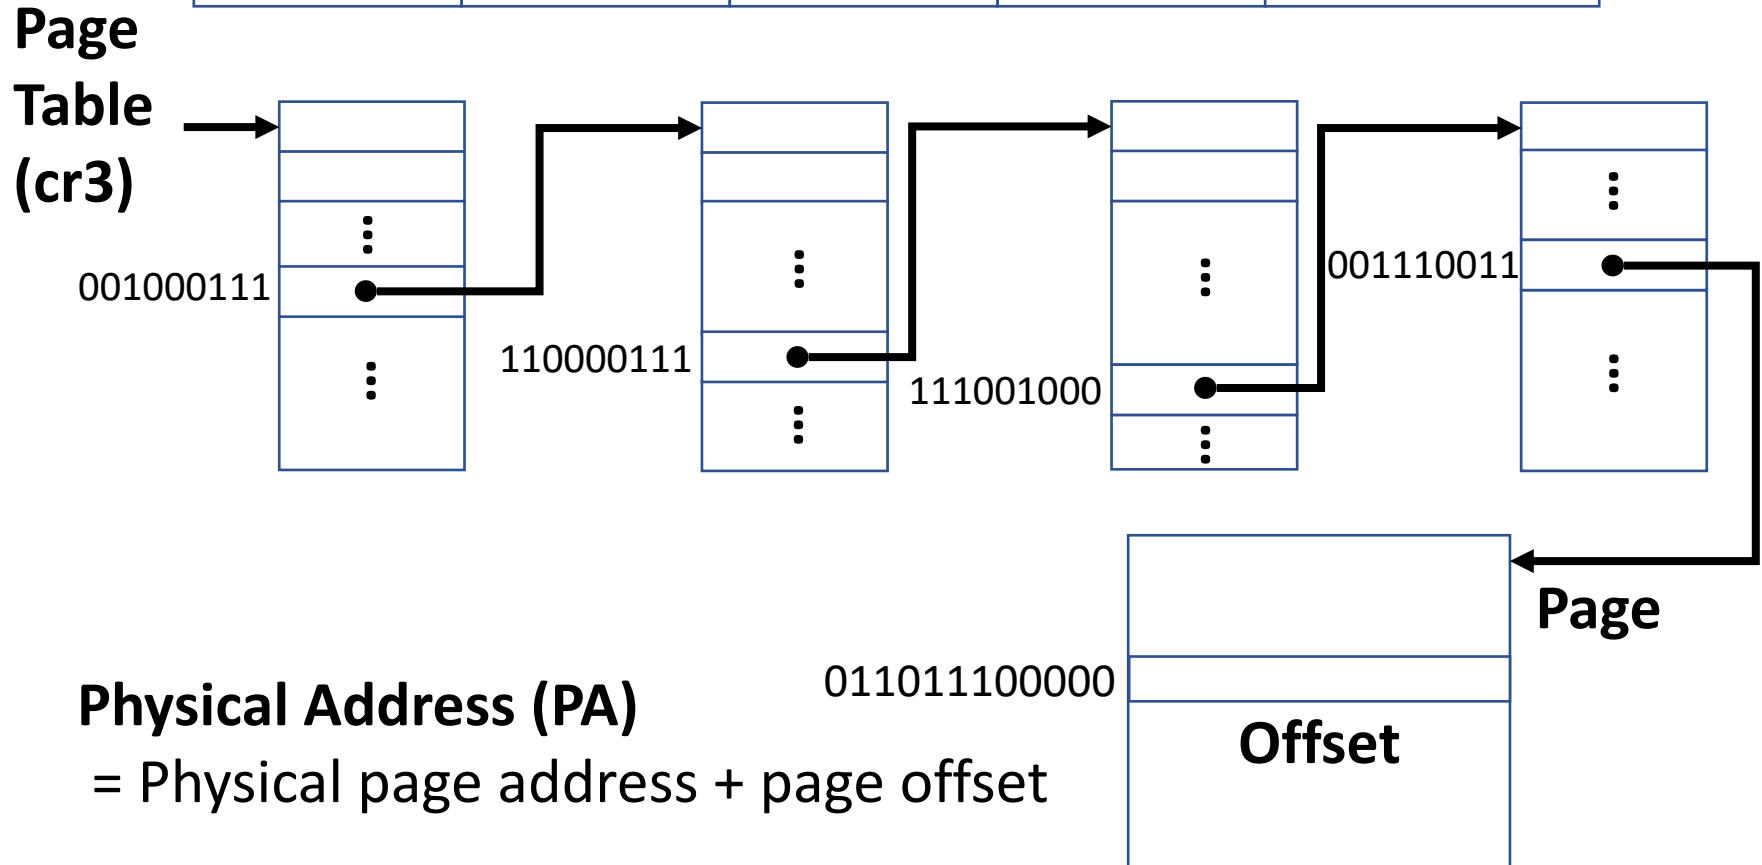
Example: Memory Management



Virtual Address → Physical Address

Virtual address (48 bits in binary)

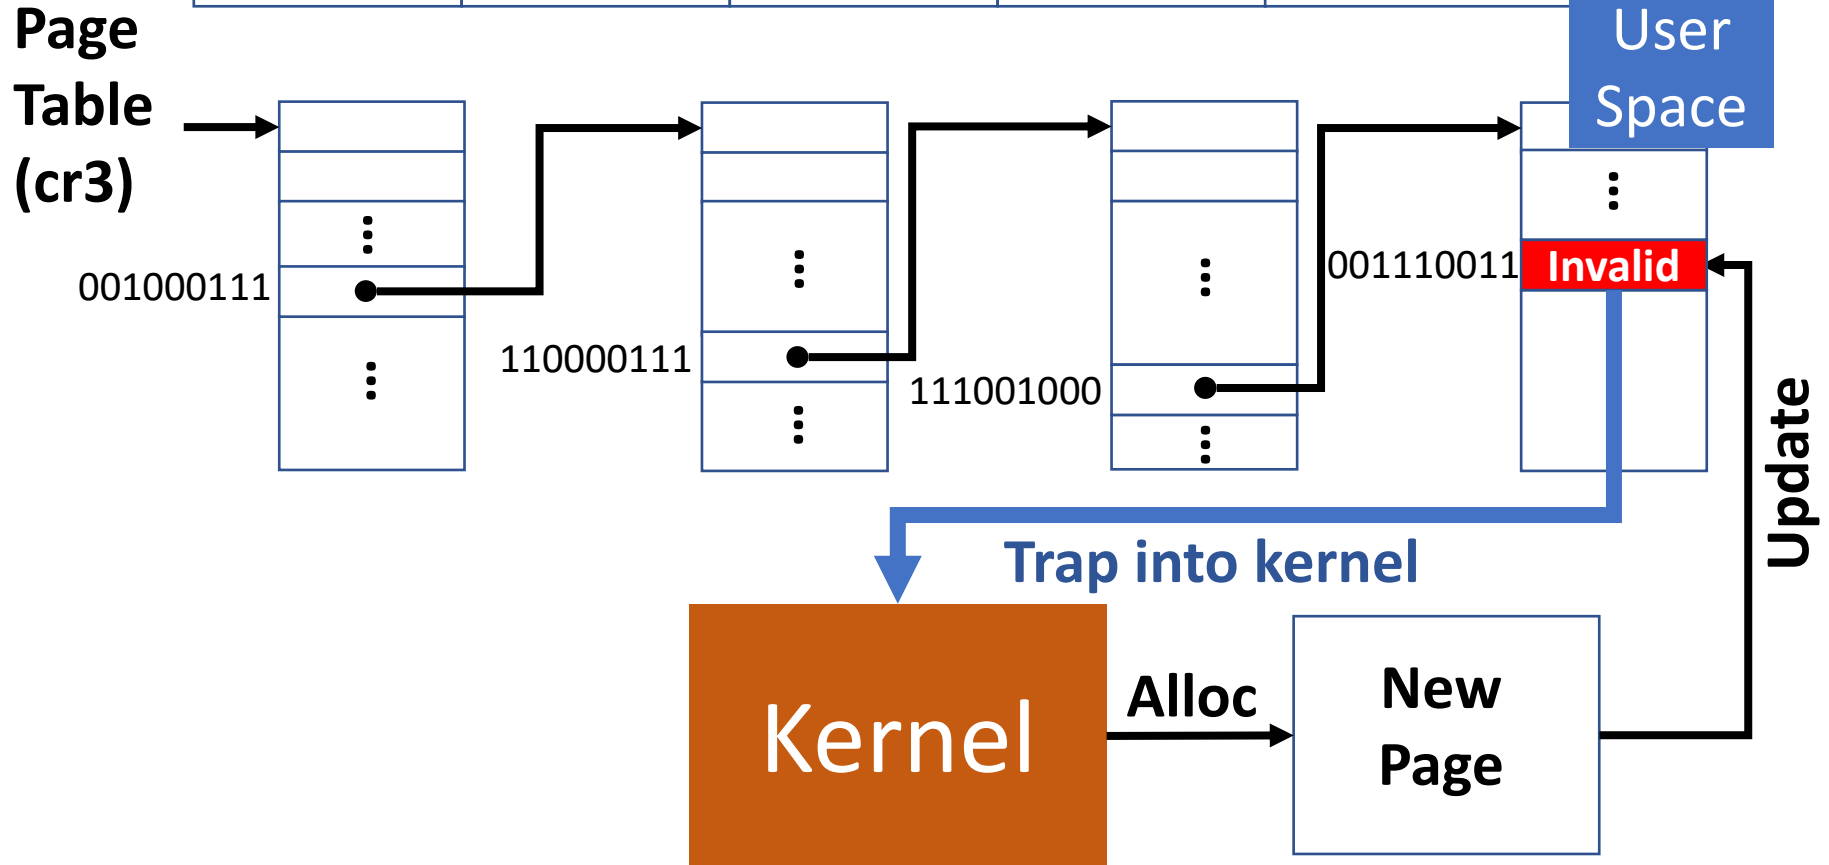
001000111	110000111	111001000	001110011	011011100000
-----------	-----------	-----------	-----------	--------------



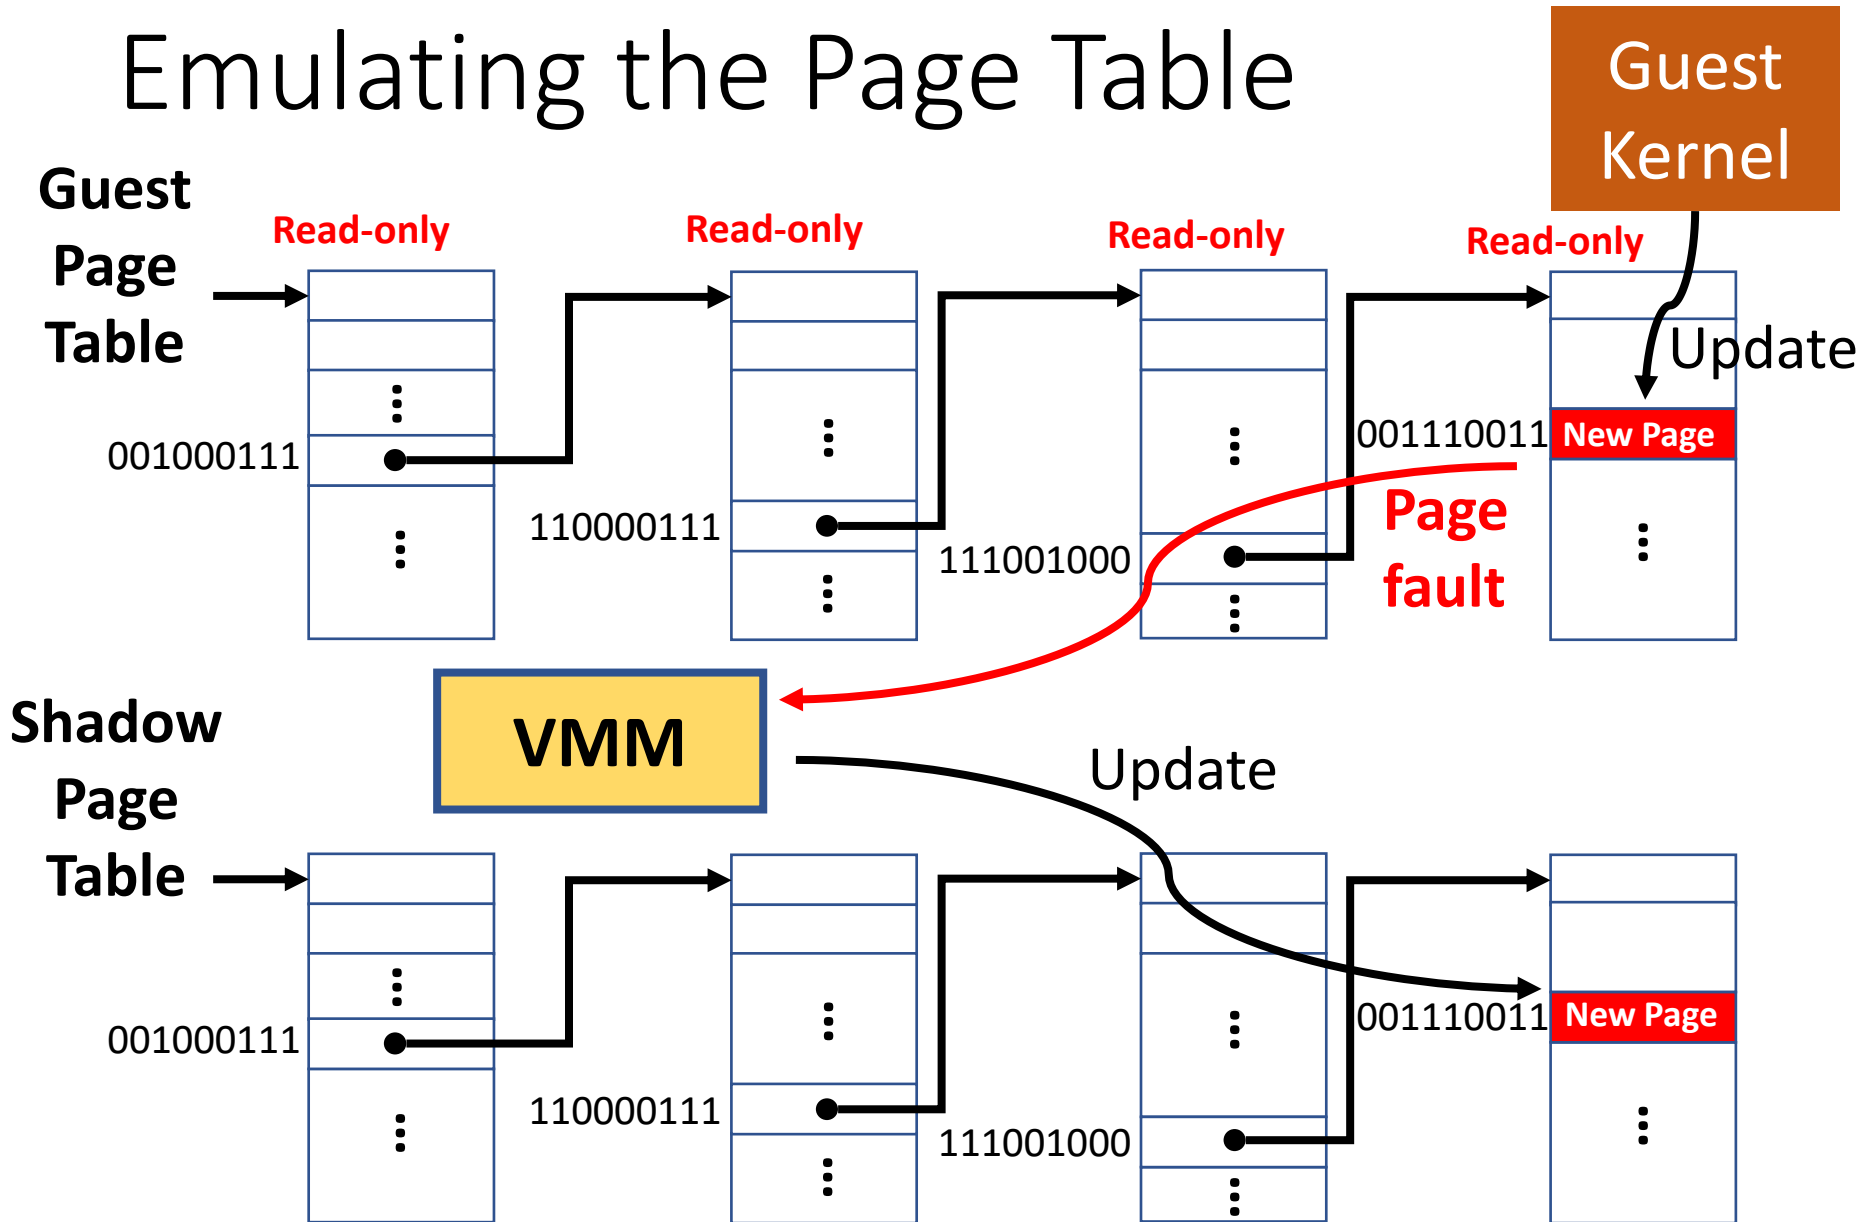
Page Fault Handling

Virtual address (48 bits in binary)

001000111	110000111	111001000	001110011	011011100000
-----------	-----------	-----------	-----------	--------------



Emulating the Page Table



Kernel vs VMM

- An OS relies on hardware protections to ensure all resources are controlled by kernel
- VMM has the exact same requirement


➔ The main challenge: How to subvert the control of OS kernel when the VMM is in charge?

Virtualization: The VMWare Approach

x86 Was Not Virtualizable

- Popek & Goldberg: all privileged instructions need to be trapped in non-kernel mode (ring > 0)
- Many x86 instructions are not trappable
 - Example 1: PUSH %cs pushes current protection level on the stack, so guest kernel can see ring != 0
 - Example 2: POPF can enable/disable interrupt in ring 0, but silently ignored in ring > 0
 - **VMM never gets the chance to emulate!**

Dynamic Binary Translation

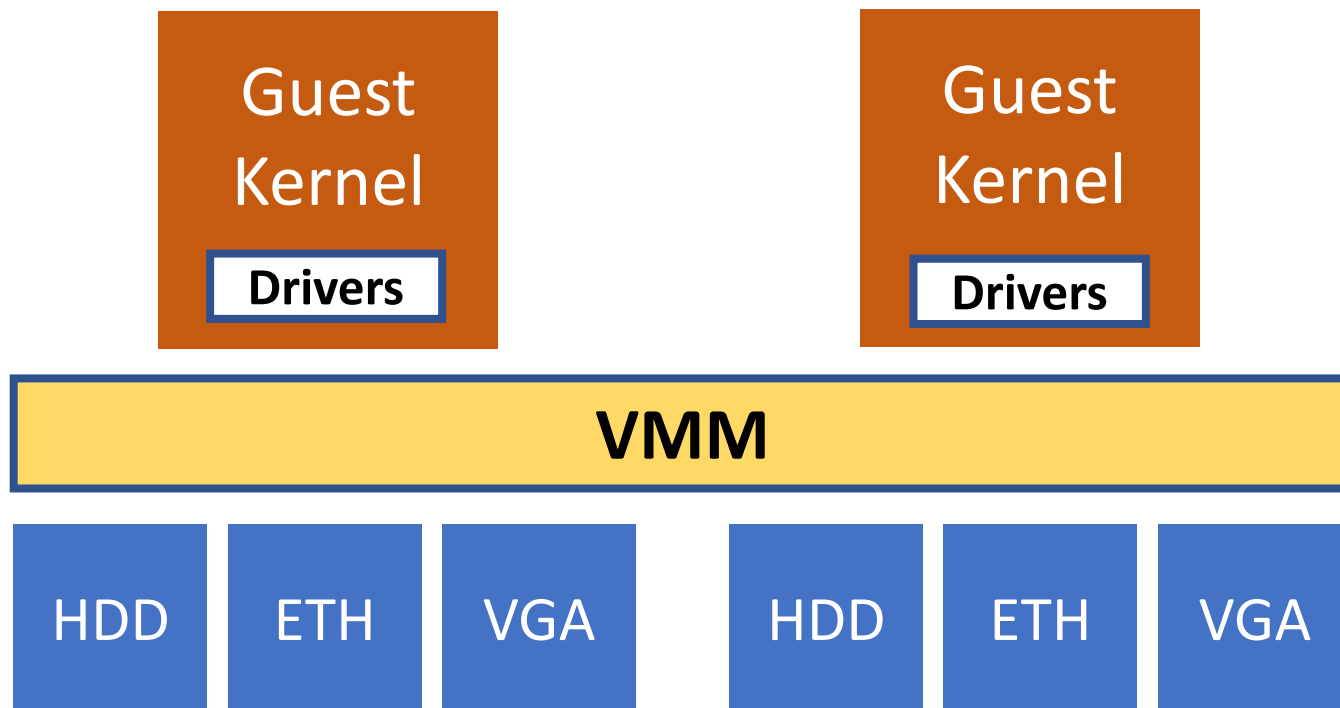


Untrapped privileged instructions can be either:

- (1) Replaced with emulation code which originally runs in VMM
- (2) Injected with other trappable instructions (e.g., syscall)

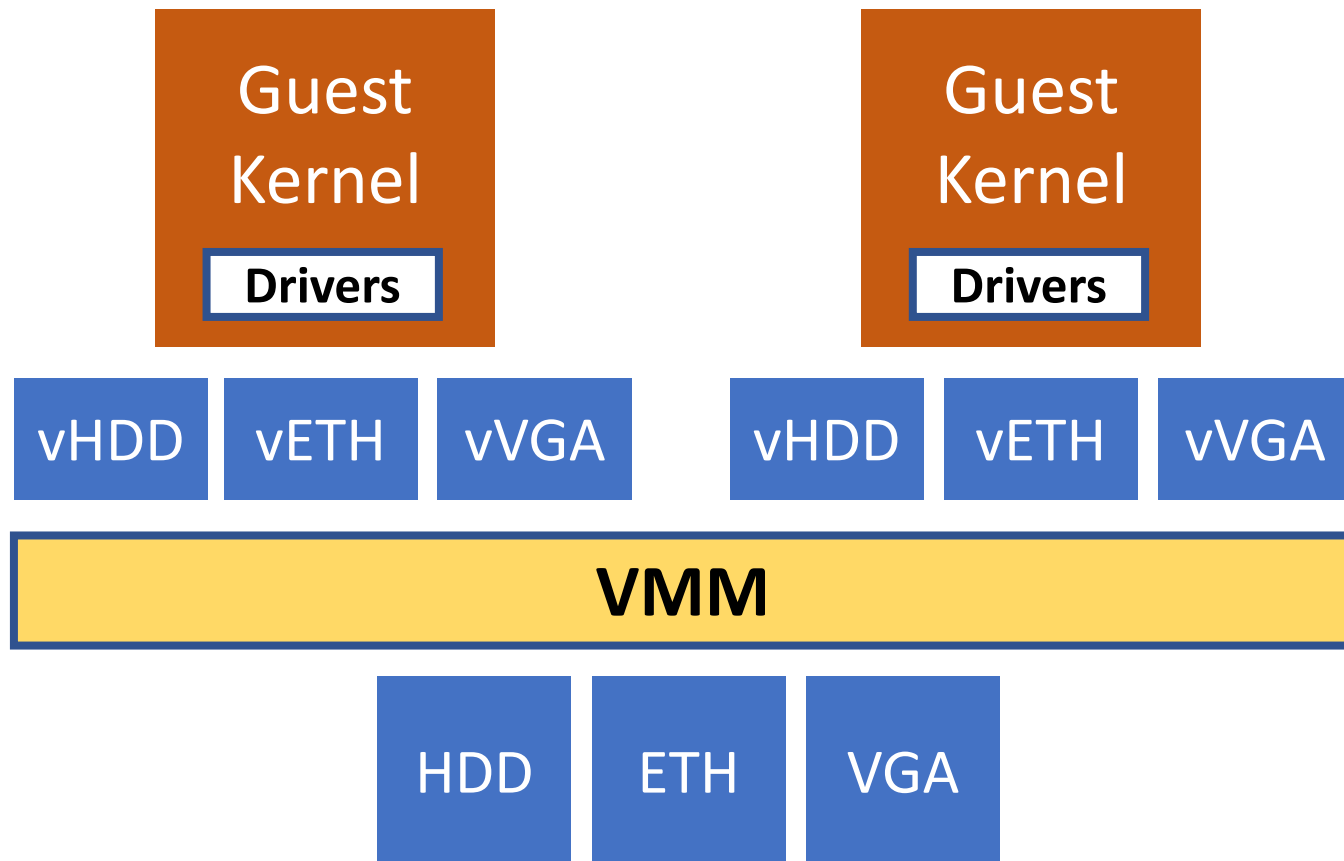
Virtualizing I/O Devices (1/2)

- Assigning physical devices to guest kernels pose engineering and security challenges because guest drivers can't access devices directly

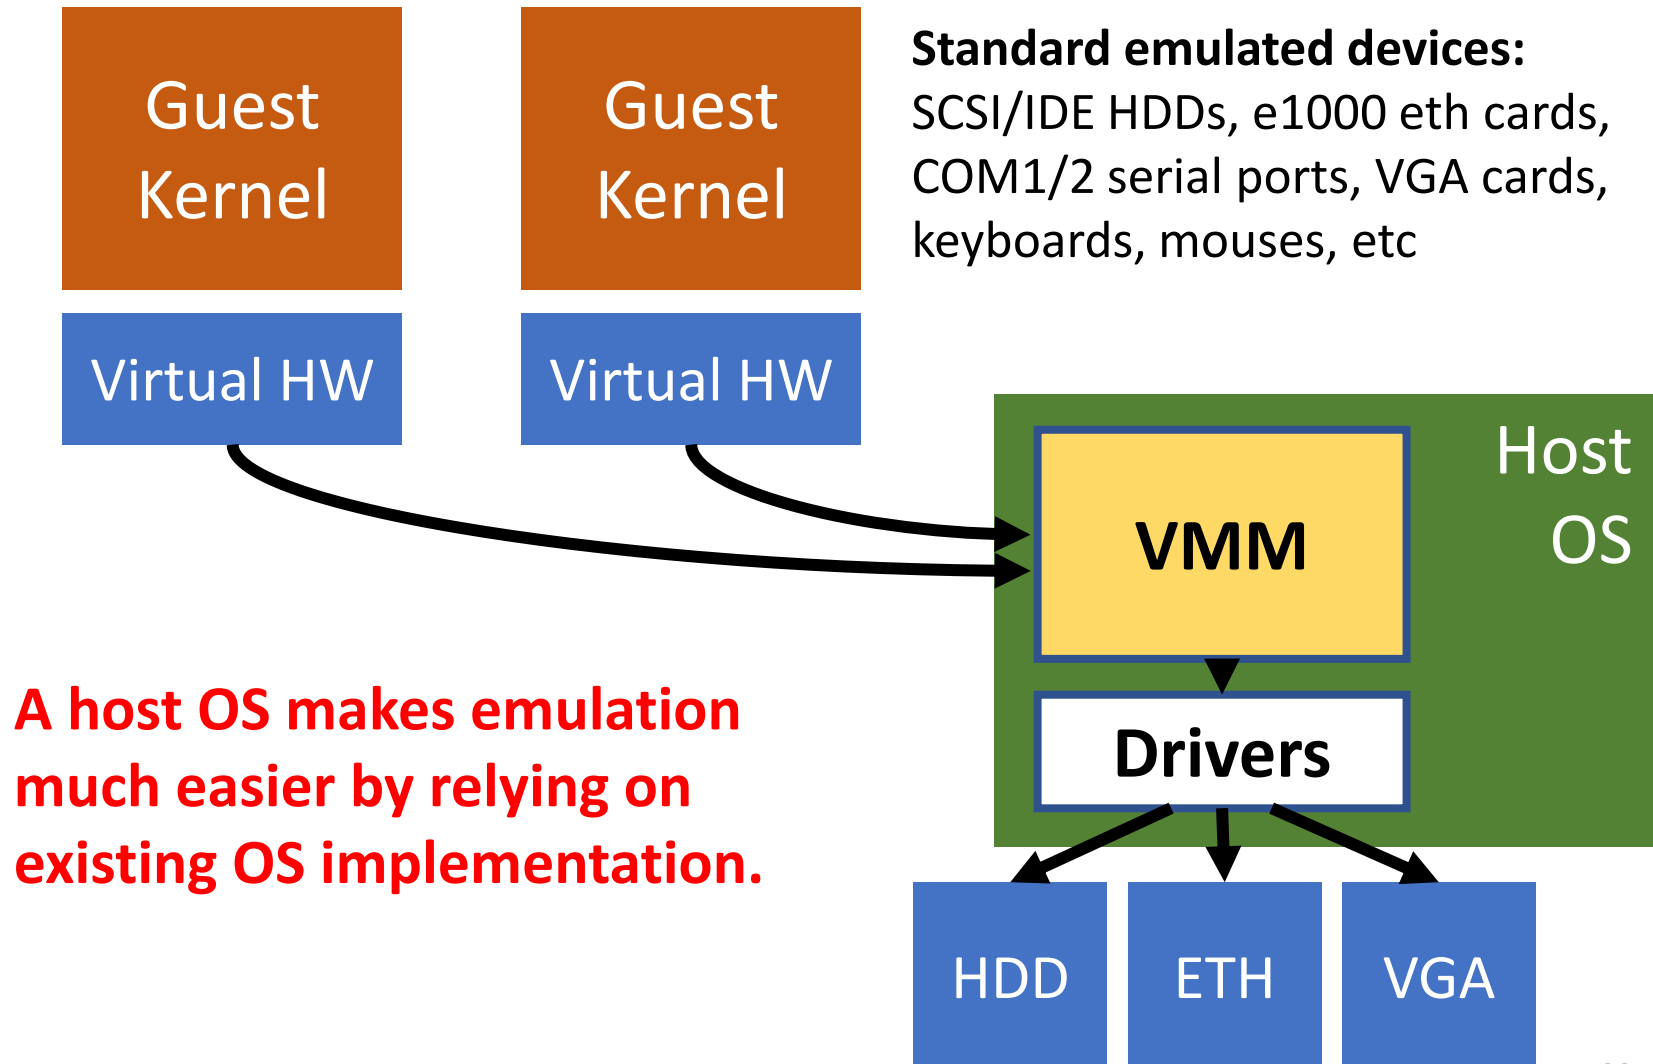


Virtualizing I/O Devices (1/2)

- VMM creates virtual devices to multiplex access to I/O resources



Use of A Host Operating System



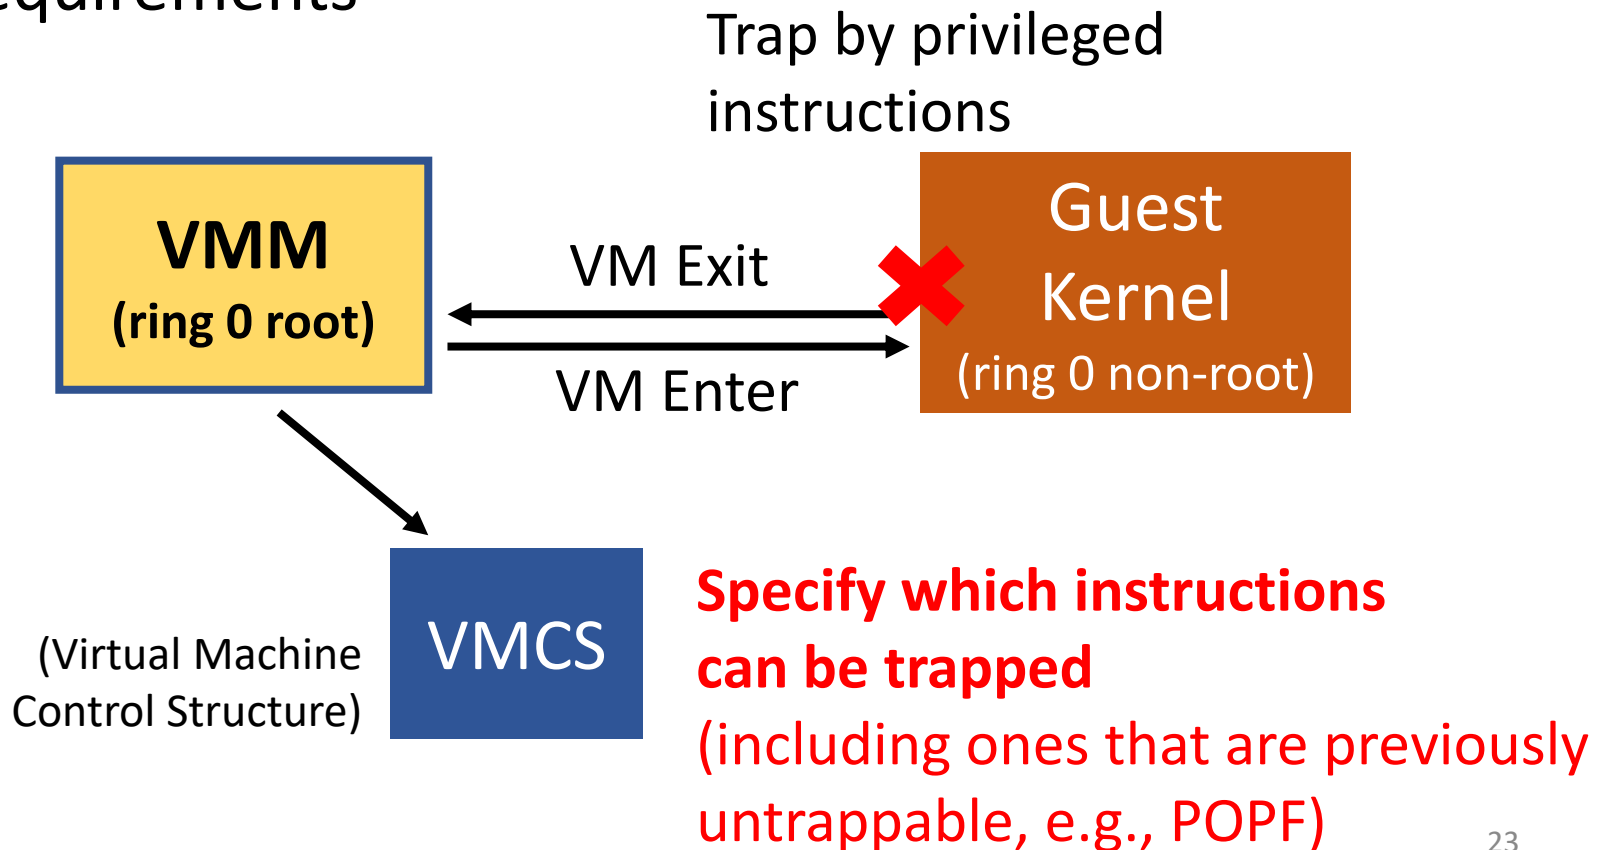
Virtualization: The Hardware Approach

Modern Virtualization Solutions

- Virtualization nowadays are assisted by hardware
 - x86 is now virtualizable
 - Hardware-assisted paging replaced shadow paging
 - Virtualization-friendly I/O devices

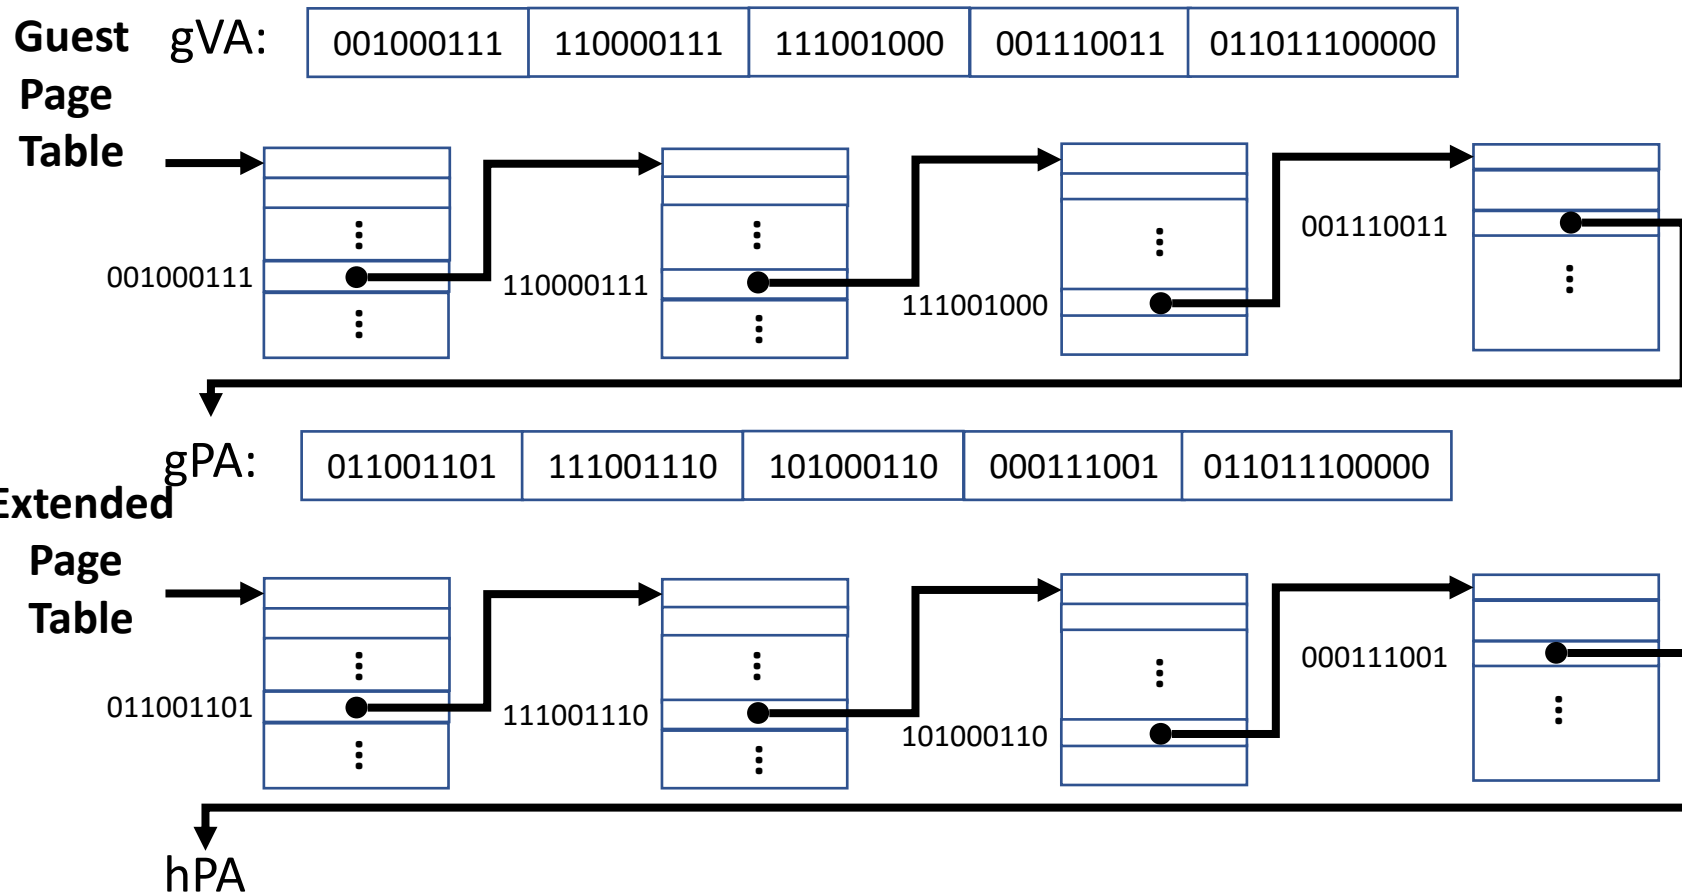
Fixing x86 Virtualization

- Intel VT or AMD-V fulfilled the Popek & Goldberg requirements



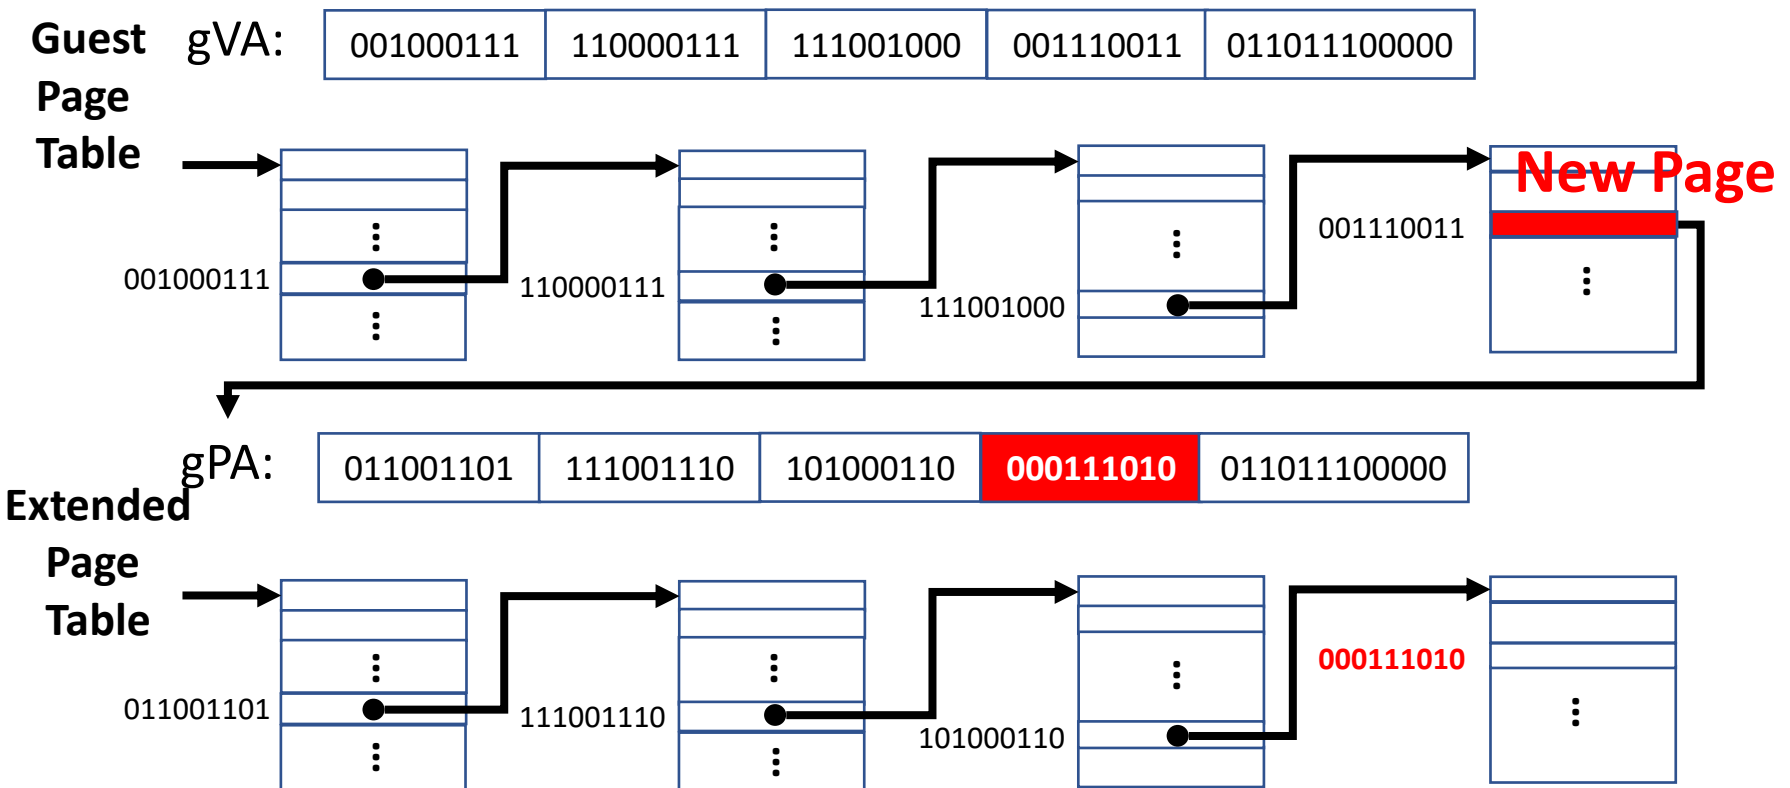
Memory Virtualization (1/3)

- Intel VT-x (extended page table) & AMD SVM (nested page table)



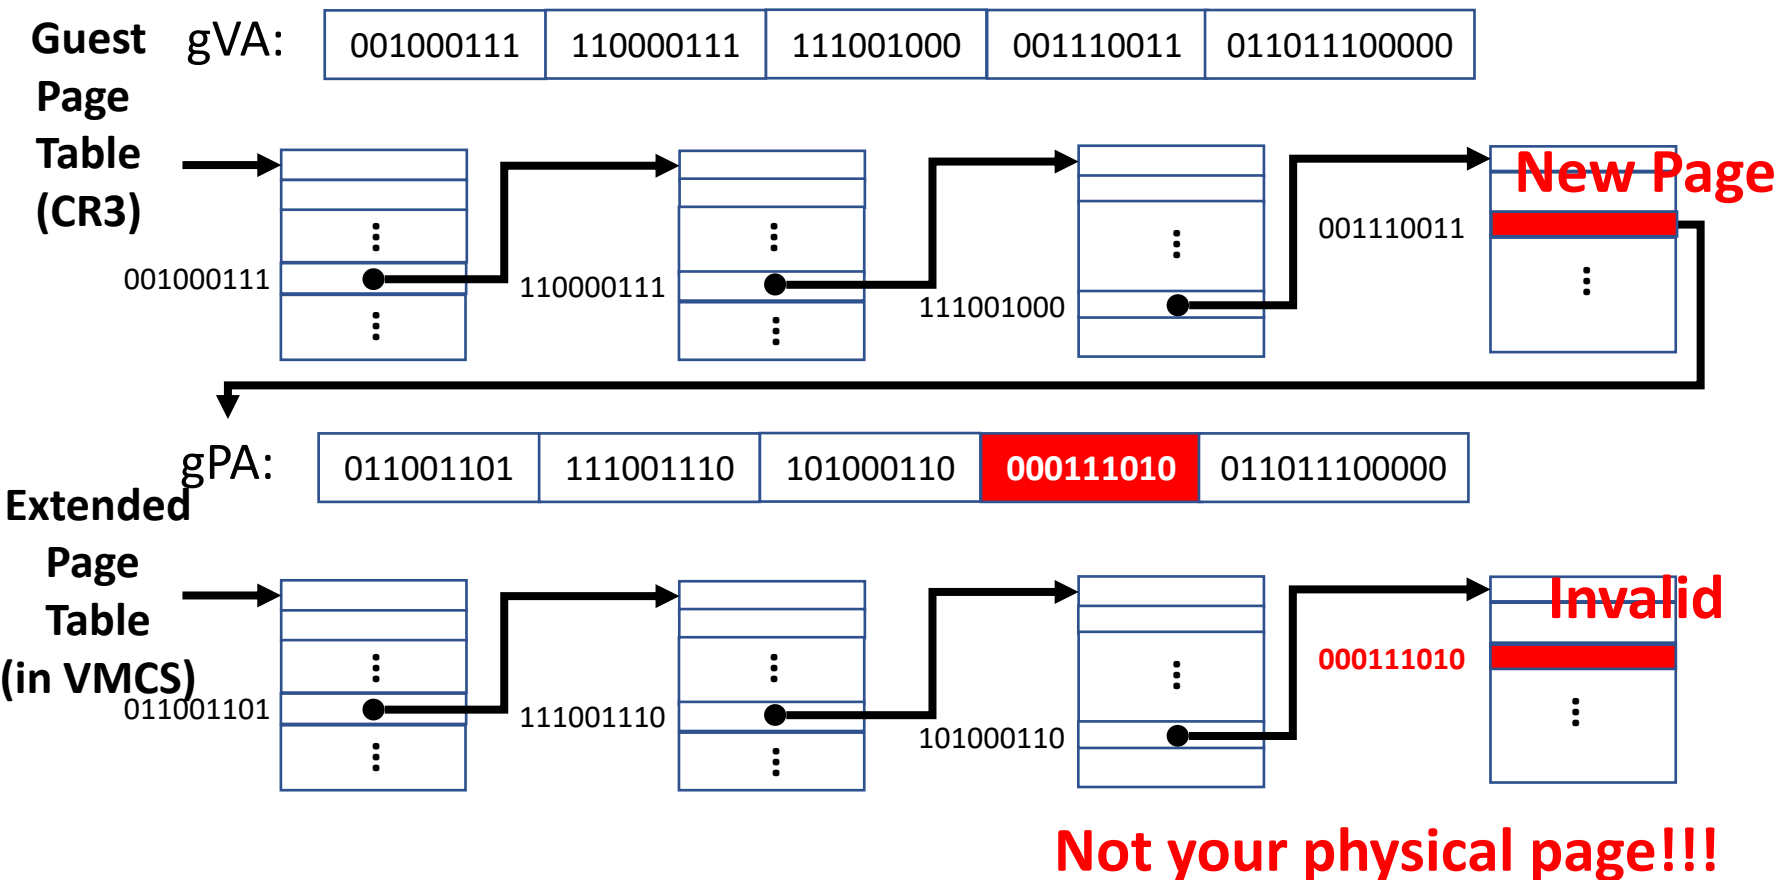
Memory Virtualization (2/3)

- Guest can update CR3 or page tables w/o trapping into VMM



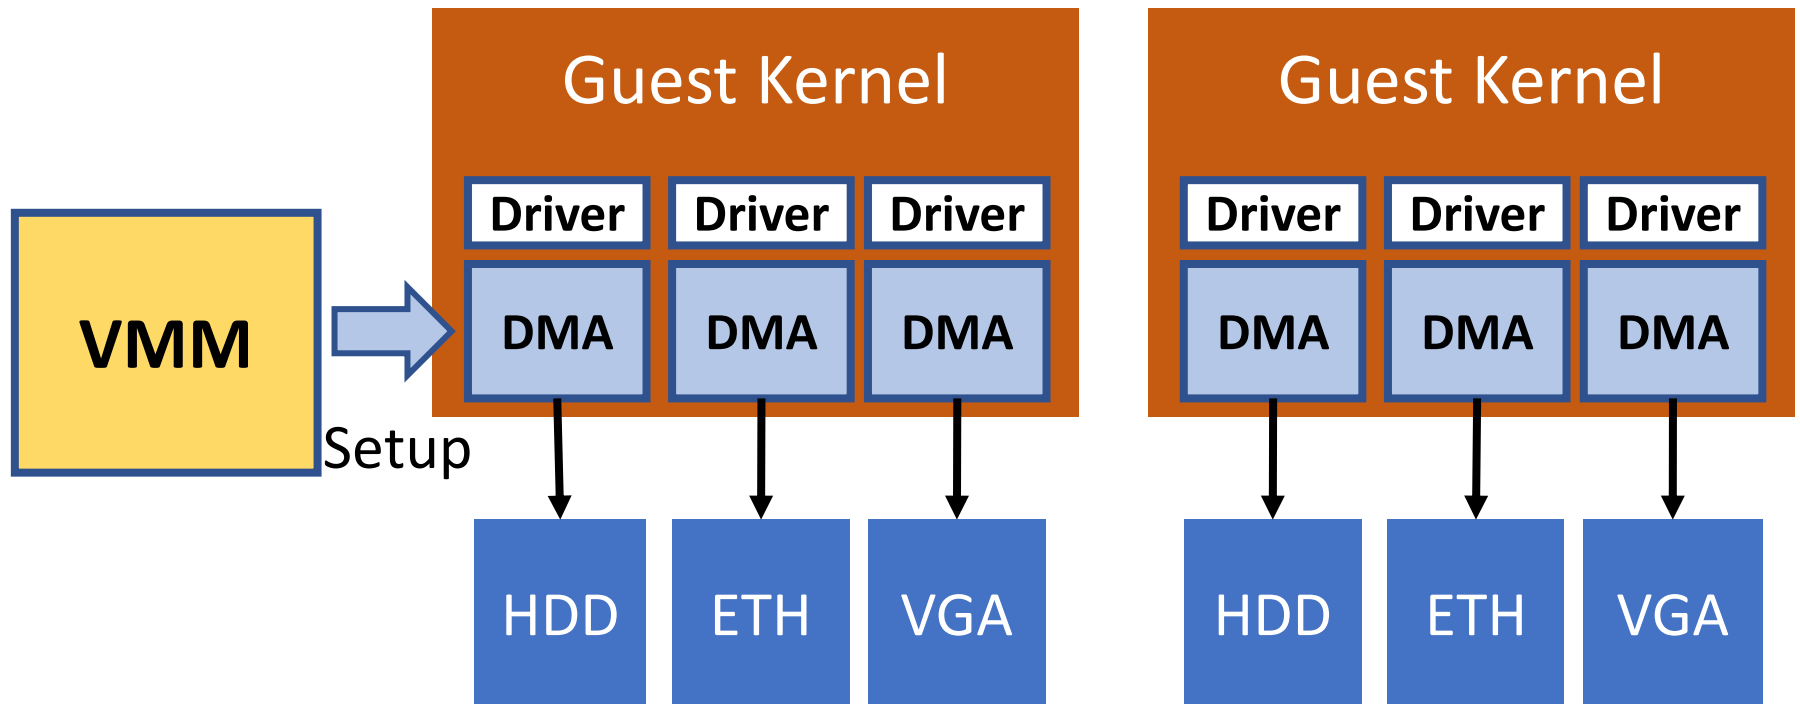
Memory Virtualization (3/3)

- VMM only manages the guest physical pages



I/O Virtualization (1/2)

- Intel VT-d and AMD-Vi allow direct I/O to assigned devices
→ Virtual IOMMU (vIOMMU)

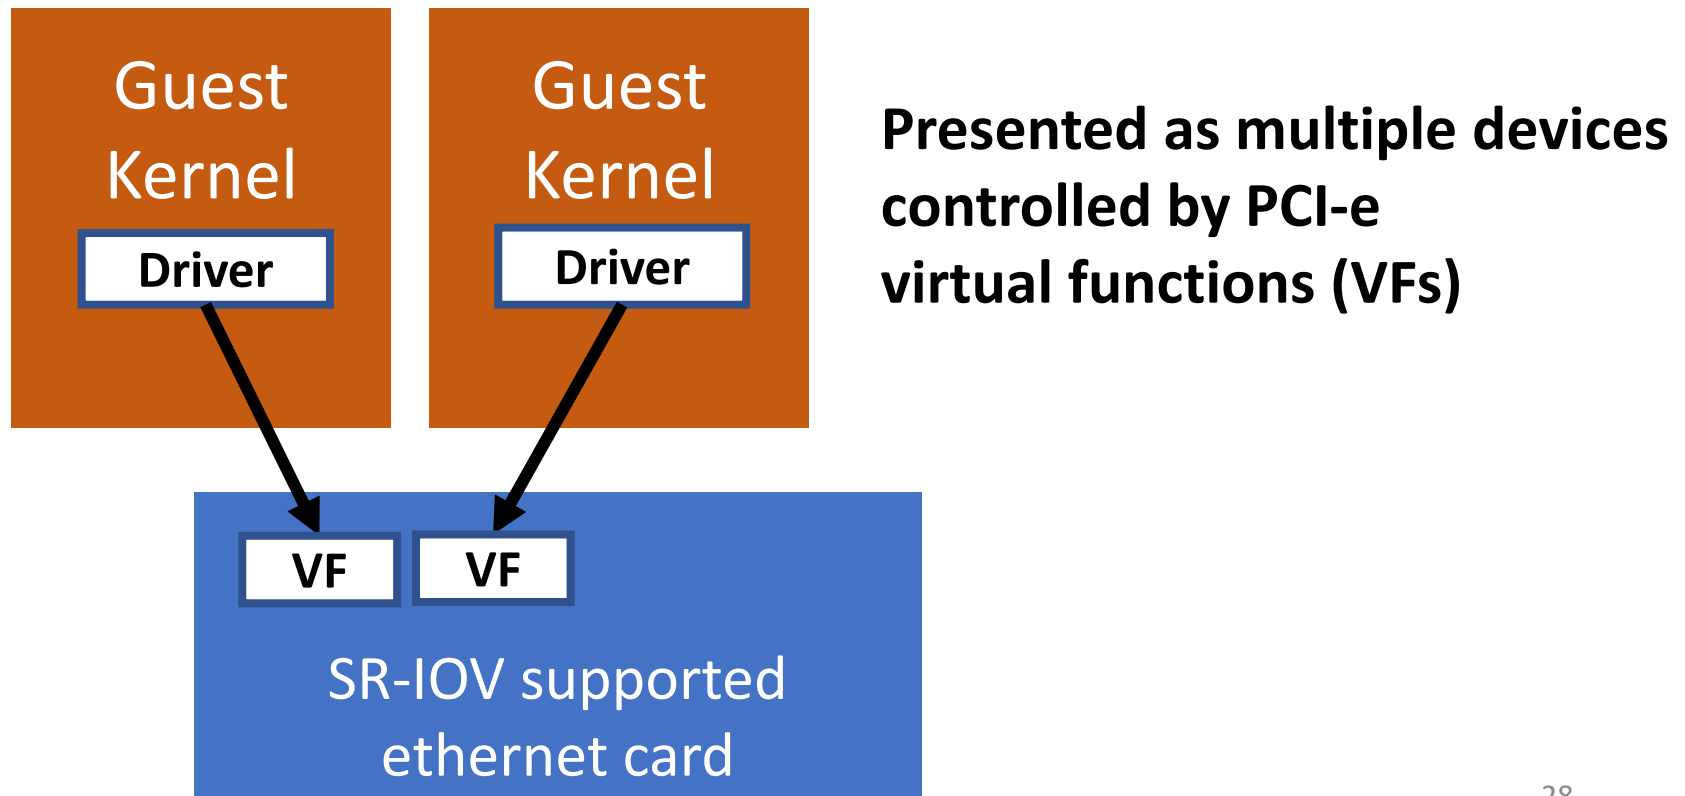


Still can't share devices!

I/O Virtualization (2/2)

- **Single-Root I/O Virtualization (SR-IOV):**

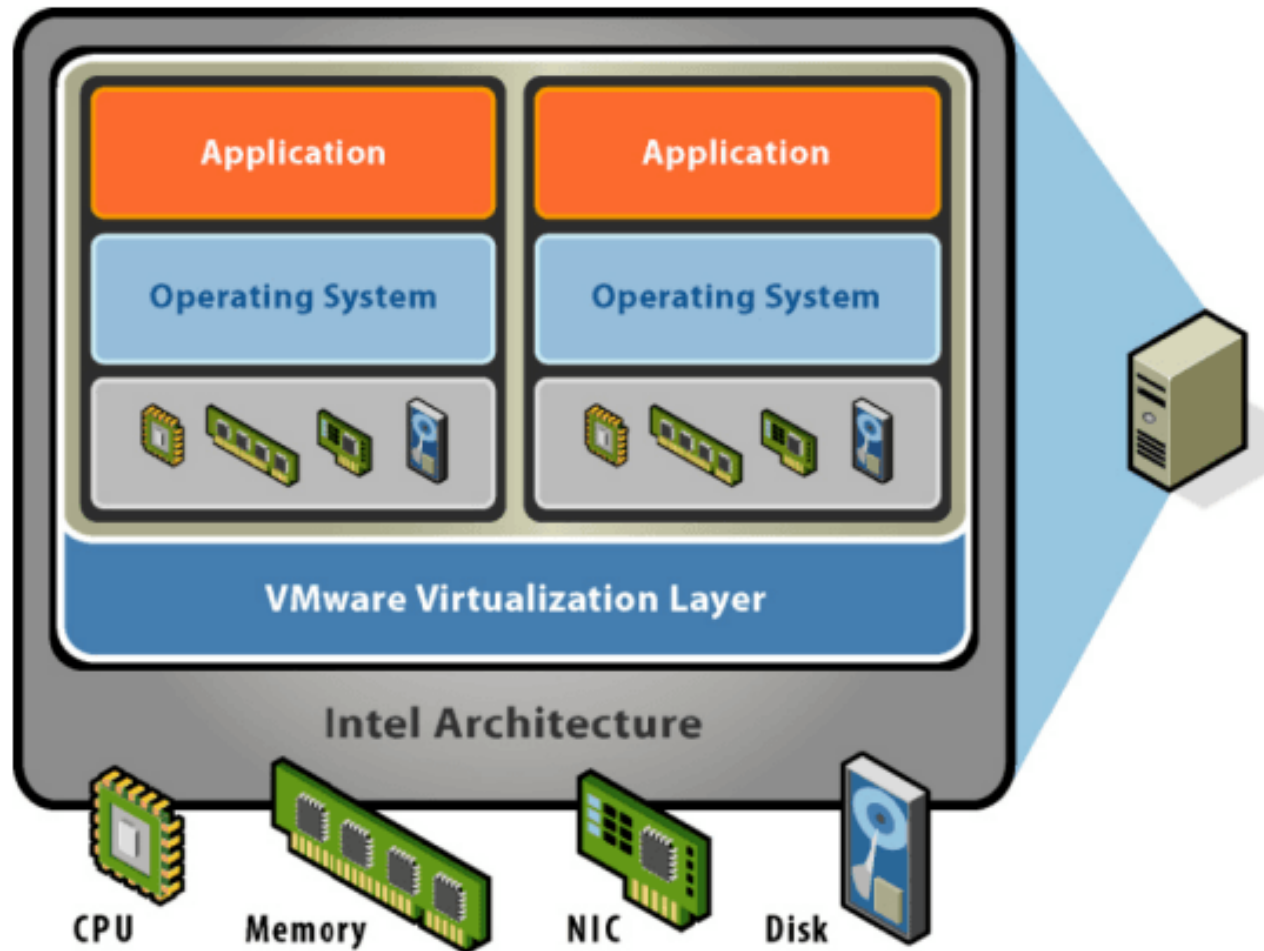
A specification for sharing PCI-e devices with multiple guests



Modern Hypervisors

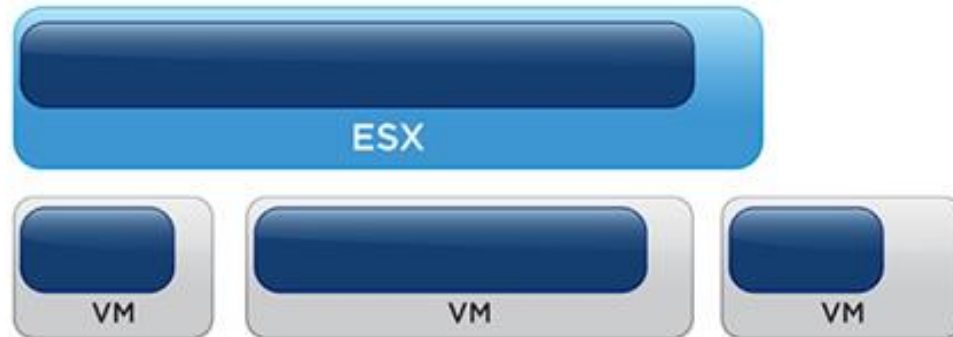
VMWare ESX/ESXi

“Bare-metal Hypervisor”

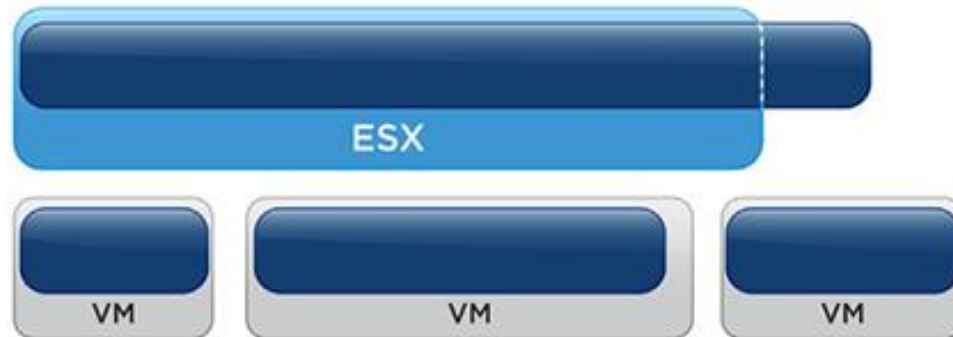


Memory Overcommitment

Not overcommitted:



Overcommitted:




 Mapped Memory

Memory Sharing

Same paging merging:


If VM1 and VM2 contain pages with exactly same contents, merge them into one physical page.




Ballooning

Reclaim memory from VMs

- A balloon module is loaded into the guests
- The balloon works on pinned physical pages in the VM
- “Inflating” the balloon reclaims memory




VM Orchestration




Snapshot/Checkpoint/Migration

Snapshot

- Save the state of VM inside the disk
- Easily restore later to resume the prior execution



Live Migration



VM Introspection

